

Data Structures

10. Binary Search Trees

10. Binary Search Trees


1. Binary Search Trees
2. Balanced binary search trees : AVL Tree

Binary Search Tree 이진탐색트리

- Heap
 - Provides a good performance of $O(\log_2 n)$ only when handling the root
 - For searching other item than the root, it takes $O(n)$ time
- Why binary search tree is necessary
 - requires $O(h)$ for any item, which is proportional to tree height (h)


Binary Search Tree

- Assume “no overlapping elements” in Tree
- The root $>$ left sub-tree nodes
- The root $<$ right sub-tree nodes
- Left & right sub-trees also must be binary search tree




Binary Search Tree


- left sub-tree < node < right sub-tree



(0)




(0)



(X)

Binary Search Tree

- Search a key in BST
 - Recursive & iterative versions



```
tree_ptr rBST(tree_ptr root, int key)
{
    if (!root) return NULL;
    if (key == root->data) return root;
    if (key < root->data)
        return rBST(root->left, key);

    return rBST(root->right, key);
}
```

$O(\log_2 n) + \text{recursion overhead}$


```
tree_ptr iBST (tree_ptr tree, int key) {
    while (tree) {
        if (key == tree->data)
            return tree;
        if (key < tree->data)
            tree = tree->left;
        else
            tree = tree->right;
    }
    return NULL;
}
```

$O(\log_2 n)$

Binary Search Tree

- Insertion Algorithm

- Before inserting a node, it should confirm whether there is no overlapping node. If overlapped, insertion fails
- Otherwise, attach a new node (35) to the last item compared (40)



Binary Search Tree

• Overlap_check()


- If tree is empty, or any overlapped node exists, NULL is returned
- Otherwise, a node searched lastly will be returned

• Time complexity : $O(\underline{\log_2 N})$

- search : $O(\underline{\log_2 N})$
- insert : $O(\underline{1})$

Binary Search Tree

```
void insert_node(tree_ptr *node, int num)
{
    tree_ptr ptr, temp = overlap_check(*node, num); //lastly found node
    if (temp || !(*node)) { // no overlap, or empty tree
        ptr = (tree_ptr) malloc (sizeof(node));
        if (IS_FULL(ptr)) {
            printf("The memory is full \n");    exit(1);
        }
        ptr->data = num;
        ptr->left = ptr->right = NULL;
        if (*node) {
            if(num < temp->data)
                temp->left = ptr;
            else
                temp->right = ptr;
        } else // empty tree
            *node = ptr;
    }
}
```




Binary Search Tree

```
tree_ptr overlap_check(tree_ptr tree, int key)
{
    tree_ptr temp = tree;

    while (tree)
    {
        temp = tree;
        if (key == tree->data) //overlapped
            return NULL;
        if (key < tree->data)
            tree = tree->left_child;
        else
            tree = tree->right_child;
    }

    return temp;
}
```



Binary Search Tree

- Deletion Algorithm

- ① deleting a leaf node

- Assign NULL to the parent's link

- ② deleting a node with one child

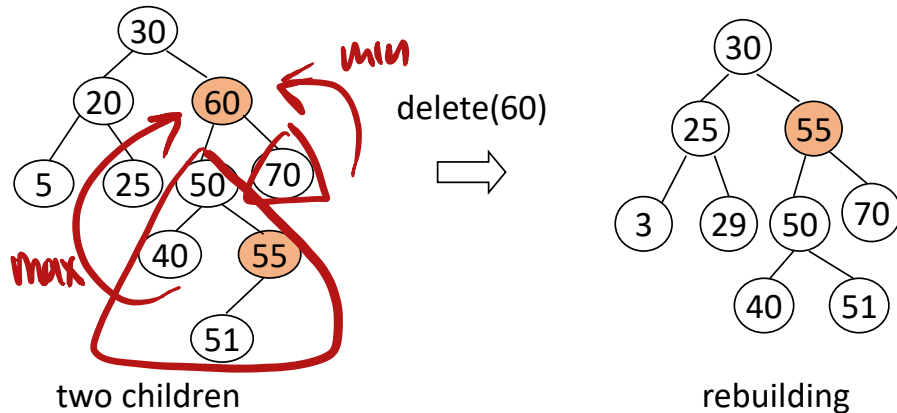

- Assign its child to the parent's link

- ③ deleting a non-leaf node with two children

- choose either the max node in left sub-tree or the min node in right sub-tree which can reduce the tree height
 - substitutes chosen node for the node to be deleted
 - Rebuild the sub-tree for BST


- Performance => $O(\log_2 n)$

Binary Search Tree



Binary Search Tree

- Time complexity
- Average case : $O(\log_2 n)$
- Worst case : skewed tree $\Rightarrow O(\underline{n})$
- Best case : balanced binary search tree $\Rightarrow O(\log_2 n)$



Balance Binary Search Tree 균형이진탐색트리


- Height of binary search tree : n
 - Insertion, deletion can be $O(n)$ in the worst case
- Good to keep a tree height small
- Minimum height of a binary tree with n nodes : $O(\log_2 n)$
- Goal
 - keep the height of a binary search tree $O(\log_2 n)$
- Balanced binary search trees
 - AVL tree, 2-3-4 tree, red-black tree

Balanced Tree?


- Suggestion
 - every node must have left and right subtrees of the same height
 - Hard to satisfy this except for complete full binary trees
- Our choice
 - for each node, the height of the left and right subtrees can differ at most 1

AVL Tree

- Adelson-Velskii and Landis, 1962
- AVL tree is a binary search tree in which
 - for every node in the tree, the height of the left and right sub-trees differ by at most 1



AVL tree




AVL property violated at ____


Balance factor

- Balance factor (BF) of a node
= Height (left subtree of the node) – Height (right subtree)
- AVL tree : BF of all node should be 1, 0 or -1

균형인수



AVL tree



AVL property
violated here


AVL Tree with Min Number of Nodes

$N_h =$ MINIMUM ^{AVL ✓} # of nodes in a tree of height h




$$N_1 = 1$$


$$N_0 = 0$$



$$N_2 = 2$$




$$N_3 = 4$$



$$N_4 = N_3 + N_2 + 1 = 7$$

$$N_h = N_{h-1} + N_{h-2} + 1$$

Thus, searching on an AVL tree will take $O(\underline{\log_2 n})$ time




Smallest AVL tree of height 9 N_9


$$N_h = N_{h-1} + N_{h-2} + 1$$

Insertion in AVL Tree


- Basically follows insertion strategy of binary search tree
 - But may cause violation of AVL tree property
- Restore the destroyed balance condition if needed



Original AVL tree




Insert 6
Property violated




Restore AVL property

Insertion in AVL Tree

- After an insertion, only nodes that are on the path from the insertion point to the root might have their balance altered
 - Because only those nodes have their subtrees altered
- Rebalance the tree at the deepest, such node guarantees that the entire tree satisfies the AVL property



Node 5, 8, 7 might
have balance altered



Rebalance node 7
guarantees the whole tree be AVL

Cases for Rebalance

- Denote the node that must be rebalanced α
 1. _____ Case : an insertion into the left subtree of the left child of α
 2. _____ Case : an insertion into the right subtree of the left child of α
 3. _____ Case : an insertion into the left subtree of the right child of α
 4. _____ Case : an insertion into the right subtree of the right child of α


Rebalance of AVL Tree

- Rebalance of AVL tree are done with simple modification to tree, known as rotation
- Insertion occurs on the “outside”
 - left-left or right-right cases
 - Is fixed by single rotation of the tree
- Insertion occurs on the “inside”
 - left-right or right-left cases
 - is fixed by double rotation of the tree

Insertion Algorithm

- First, insert a new key as a new leaf just as in ordinary binary search tree
- Check BF of each node in the path between a new node (N) and the root.
 - If BF is OK(0, -1, +1), proceed to parent(x)
 - If not, restructure it by doing either a single rotation or a double rotation
- Note
 - once we perform a rotation at a node x, we won't need to perform any rotation at any ancestor of x.



Single Rotation (Left-Left Insertion)



An insertion in subtree X ,
AVL property violated at node k_2


Solution: single rotation

Single Rotation (Left-Left insertion)




Single Rotation (Right-Right insertion)

- Case 4 is a symmetric case to case 1
- Insertion takes $O(\text{_____})$ time, Single rotation takes $O(1)$ time



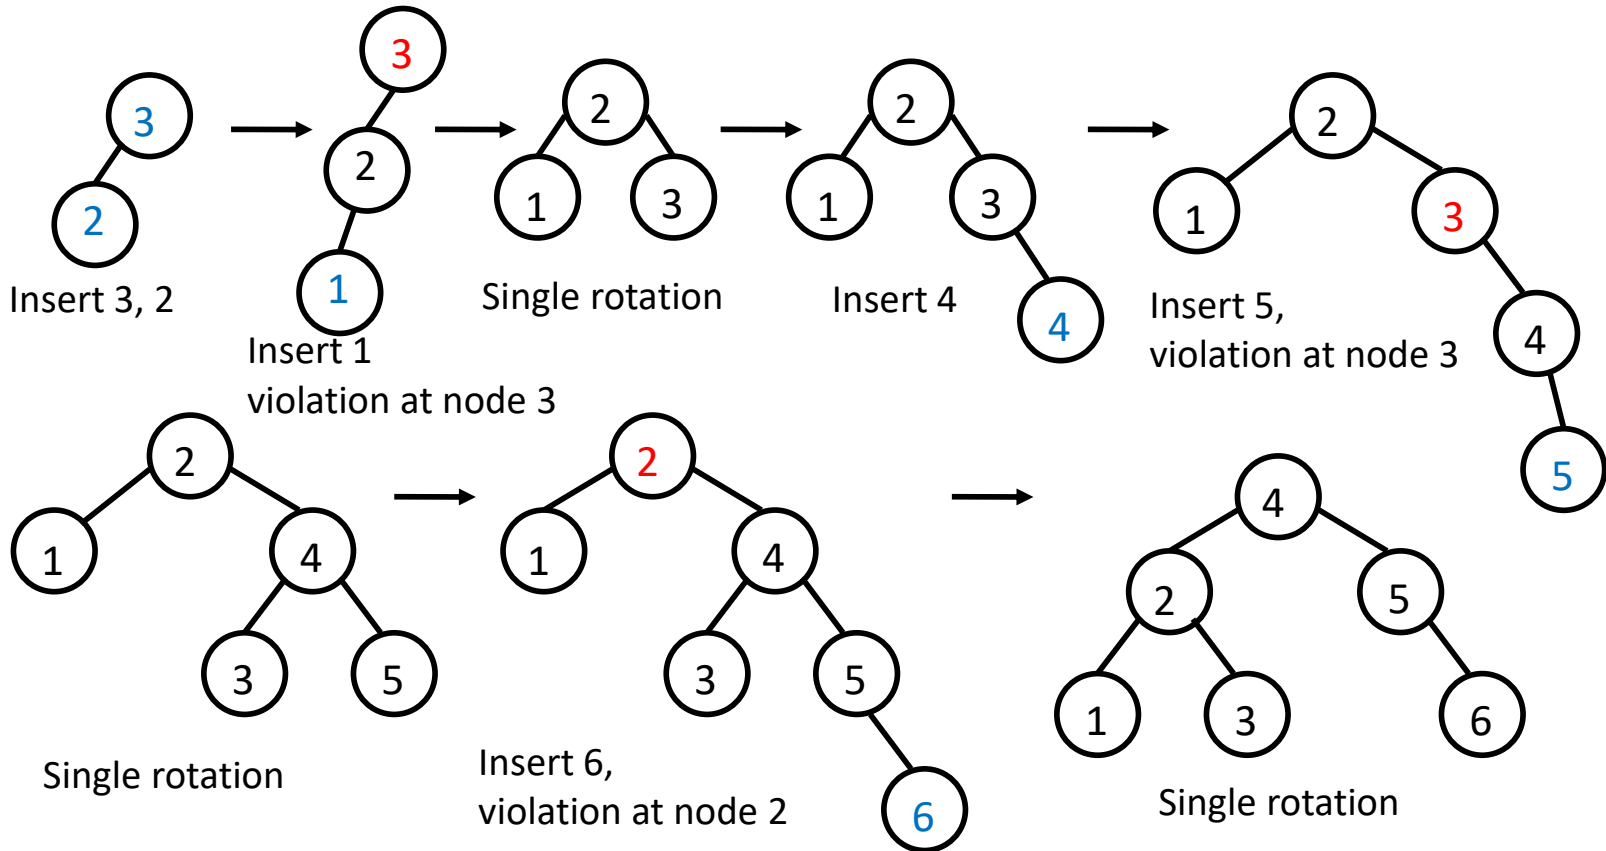
Single Rotation (Right-Right insertion)

- right-right insertion





AVL Tree Construction



- Sequentially insert 3, 2, 1, 4, 5, 6 to an AVL Tree



AVL Tree Construction



Single rotation



Single rotation,
Violation remains


Double Rotation (Left-Right insertion)

- Facts

- A new key is inserted in the subtree B or C
- The AVL-property is violated at ____
- k_3 - k_1 - k_2 forms a zig-zag shape

- Solution


- The only alternative is to place ____ as the new root



Double rotation to fix case 2

Double Rotation (Right-Left insertion)



- Facts
 - The new key is inserted in the subtree B or C
 - The AVL-property is violated at _____
 - k_1 - k_3 - k_2 forms a zig-zag shape
- Case 3 is a symmetric case to case 2




Double rotation to fix case 3


AVL Tree Construction

- continue to insert 15, 14, 13, 12, 11, 10, 8, 9







Insert 14







Double rotation





Insert 13





Single rotation





Single rotation



Double rotation

5/19(5)

B-Trees

Motivation for B-Trees

- Index structures for large datasets cannot be stored in main memory
- Storing it on disk requires different approach to efficiency
- Assuming that a disk spins at 7200 RPM, one revolution occurs in $1/120$ of a second, or 8.3ms
- Crudely speaking, one disk access takes about the same time as 200,000 instructions

Motivation (cont.)

- Assume that we use an AVL tree to store about 20 million records
- We end up with a very deep binary tree with lots of different disk accesses; $\log_2 20,000,000$ is about 24, so this takes about 0.2 seconds
- We know we can't improve on the $\log n$ lower bound on search for a binary tree
- But, the solution is to use more branches and thus reduce the height of the tree! (자식이 최대 2개로 고정돼있으므로)
 - As branching increases, depth decreases

→ 여러개의 자식을 가질수 있게하여 높이를 줄여야한다.


Comparing Trees

- Binary trees
 - Can become *unbalanced* and *lose* their good time complexity ($O(\log_2 n)$)
 - AVL trees are strict binary trees that *overcome the balance problem*
 - Heaps remain balanced but only good to get the root(max/min)
- Multi-way trees 다원트리
 - B-Trees can be *m*-way, they can have any number of children

m-way B-tree

Definition : m-way B-Tree

- DEF: A B-Tree of order m is an m -way search tree that either is empty or satisfies the following properties

- (1) The **root** node has at least two children or a leaf
- (2) All nodes other than the root node and external nodes have at most m and at least $\lceil m/2 \rceil$ children
내부 
- (3) All **external** nodes are at the same level
- (4) The number of keys is **one less** than the number of children for **non-leaf** nodes and at most $m-1$ and at least $\lceil m/2 \rceil - 1$ for **leaf** nodes

M-way B tree

- $m = 3$ (2-3 Tree) , $\text{min children} = \text{ceil}(3/2)$
 - $\text{Degree}(\text{internal}) = 2 \text{ or } 3$, $\text{Degree}(\text{root}) = 0, 2, 3$
 - # keys in a leaf node = 1, 2
- $m = 4$ (2-3-4 tree) , $\text{min children} = \text{ceil}(4/2)$
 - $\text{Degree}(\text{internal}) = 2, 3, 4$, $\text{Degree}(\text{root}) = 0, 2, 3, 4$
 - # keys in a leaf node = 1, 2, 3
- $m = 5$, $\text{min children} = \text{ceil}(5/2)$
 - $\text{Degree}(\text{internal}) = 3, 4, 5$, $\text{Degree}(\text{root}) = 0, 2, 3, 4, 5$
 - # keys in a node = 2, 3, 4

Entries in B-trees of Various orders

Order	Number of Subtrees		Number of Entries	
	Min	Max	Min	Max
3	2	3	1	2
4	2	4	1	3
5	3	5	2	4
6	3	6	2	5
...
m	$\text{ceil}(m/2)$	m	$\text{ceil}(m/2) - 1$	m-1

Creating a B-tree of order 5


- A G F B K D H M J E S I R X C L N T U P

Insert into 5-way B-tree


• A G F B K D H M J E S I R X C L N T U P



• A G F B K D H M J E S I R X C L N T U P





overflow



Split



Insert into 5-way B-tree

• A G F B K **D H M J** E S I R X C L N T U P





Insert into 5-way B-tree

• A G F B K D H M J **E S I R X** C L N T U P




Insert into 5-way B-tree

• A G F B K D H M J E S I R X **C** L N T U P




Insert into 5-way B-tree

• A G F B K D H M J E S I R X C **L N T U** P




• A G F B K D H M J E S I R X C L N T U **P**




Final

• A G F B K D H M J E S I R X C L N T U P




• Final

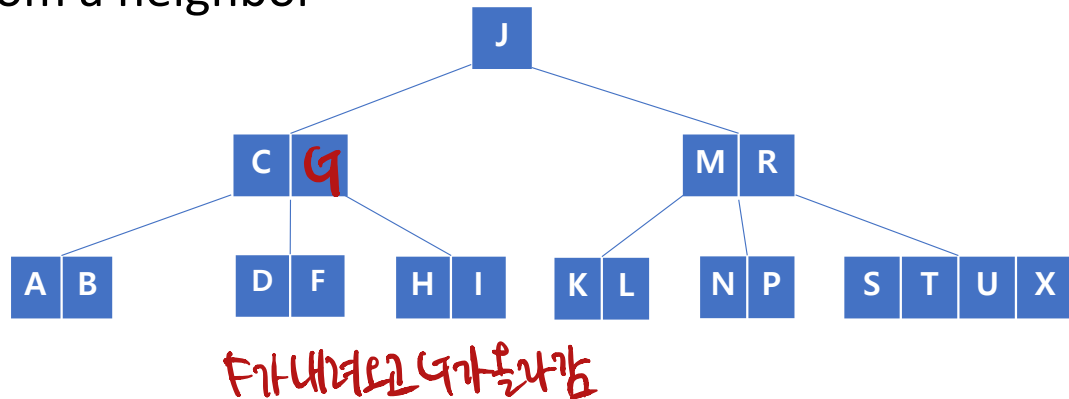


Delete

- Delete **E** from leaf node




- Borrow from a neighbor




Delete

- Delete **F** => underflow => but can't borrow from a neighbor




- Can't borrow so **combine**
- c is underflow




Delete

- so combine




- Delete **M** from non-leaf node
- Note: immediate predecessor in non-leaf is always in a leaf.
- Overwrite **M** with immediate predecessor (L) => underflow




Delete

- Borrow from a neighbor




- Delete U, X




Delete


- Delete **R**



- Underflow, can't borrow => combine



Final



Analysis of B-Trees

- The maximum number of items in a B-tree of order m and height h :

root $m - 1$

height 1 $m(m - 1)$

height 2 $m^2(m - 1)$

...

height h $m^h(m - 1)$

- So, the total number of items is

$$(1 + m + m^2 + m^3 + \dots + m^h)(m - 1) =$$

$$[(m^{h+1} - 1) / (m - 1)] (m - 1) = \underline{m^{h+1} - 1}$$

- When $m = 5$ and $h = 2$, this gives $5^3 - 1 = \underline{124}$

Reasons for using B-Trees

- When searching tables held on disc, the cost of each disc transfer is high but doesn't depend much on the amount of data transferred, especially if consecutive items are transferred
 - If we use a B-tree of order 101, say, we can transfer each node in one disc read operation
 - A B-tree of order 101 and height 3 can hold $101^4 - 1$ items (approximately 100 million) and any item can be accessed with 3 disc reads (assuming we hold the root in memory)
- If we take $m = 3$, we get a **2-3 tree**, in which non-leaf nodes have two or three children (i.e., one or two keys)
 - B-Trees are always balanced (since the leaves are all at the same level), so 2-3 trees make a good type of balanced tree