

# Transaction Processing Concepts

↳ 프로그램 독립적 작업 단위

↳ undo

transaction fail 난 경우? (transaction 종도에) ⇒ 업던 것으로 되나해야함!

read/write 후 commit을 call해야 하나의 transaction이 끝남

log에 commit이 없는 경우 transaction fail → recovery 작업 필요

↳ rollback

# Single-User VS Multi-user Systems

- Single-user vs multi-user
  - A DBMS is single-user if at most one user at a time can use the DBMS.
  - A DBMS is multi-user if many users can use the DBMS simultaneously.
- Interleaved processing vs parallel processing



# Transactions < read write


- Concurrent execution of user programs is essential for good DBMS!
- 동시성 제공 의미  
{ Usability: Multiple users or applications need to (can) access databases simultaneously.  
Performance : Since disk access speed is relatively slow, it is important to keep the cpu humming by working on several user programs concurrently.
- A transaction is the DBMS's abstract view of a user program: a logical unit of database processing that includes the sequence of reads and writes.

# Transactions

## □ Conventional notations

- $T_1, T_2, T_n$  : transactions
- $R(X)$  : Reads a database item named  $X$  into a variable  $X$ .
- $W(X)$  : Write the value of  $X$  into the database item named  $X$ .
- Read-set( $T_1$ ) : set of items that the transaction  $T_1$  reads.
- Write-set( $T_1$ ) : set of items that the transaction  $T_1$  writes.

11/21



transaction  
: sequence of operations project unit

중요. 자주

# Transactions

## □ Conventional notations Cont'd

- Example:

예제  
□ T1 is a transaction that transfers the balance (N) from an account (X) to the other account (Y)

중요  
□ T2 is a transaction that deposits the amount (M) into an account (X).

operation이 아님

읽기

X의 값

T<sub>1</sub>: BEGIN R(X), X=X-N, W(X), R(Y), Y=Y+N, W(Y) END

T<sub>2</sub>: BEGIN R(X), X=X+M, W(X) END

• Read-set(T1)? Write-set(T1) ? Read-set(T2)? Write-set(T2) ?

Read-set(T1) = {X, Y}, Write-set(T1) = {X, Y}

Read-set(T2) = {X}, Write-set(T2) = {X}

순차적 serial : 하나씩만 다음 실행

↔ 동시 실행 interleaving

# Concurrency in a DBMS

- Users submit transactions, and can think of each transaction as executing by itself.
  - Concurrency is achieved by the DBMS, which interleaves actions (reads/writes of DB objects) of various transactions.
  - Each transaction must leave the database in a consistent state if the DB is consistent when the transaction begins.
    - DBMS will enforce <sup>integrity constraints</sup> ICs, depending on the ICs declared in CREATE TABLE statements.
    - Beyond this, the DBMS does not really understand the semantics of the data. (e.g., it does not understand how the interest on a bank account is computed).
- Issues: Effect of *interleaving* transactions, and *crashes*.

# *Anomalies with Interleaved Execution*

## □ The Lost Update Problem *교신계업어짐!*

T1: BEGIN R(X),  $X = X - N$ , W(X), R(Y),  $Y = Y + N$ , W(Y) END

T2: BEGIN R(X),  $X = X + M$ , W(X) END

time

- Let's say that  $X = 100$ ,  $Y = 0$  at the start and  $N = M = 50$ . Then after processing T1 and T2, what are the values of X and Y ?
  - X should be 100.
  - Y should be 50.

# Anomalies with Interleaved Execution (Cont'd)

## □ The Lost Update Problem (Cont'd)

T1: R(X), X=X-N,	W(X), R(Y),	Y=Y+N, W(Y)
T2: R(X), X=X+M,	W(X)	

\_\_\_\_\_time

Let's say that  $X = 100$ ,  $Y=0$  at the start and  $N = M = 50$ .

X :

50

150

Y:

100 이어야 됨!


50



# Anomalies with Interleaved Execution (Cont'd)

## □ The Lost Update Problem (Cont'd)

X has incorrect value because  
its update by T1 is “lost”.



Let's say that  $X = 100$ ,  $Y=0$  at the start and  $N = M = 50$ .

X :	50	150	
Y :			50

# Anomalies with Interleaved Execution (Cont'd)

아차! 자칫하면 안될 상황

cf) dirty bit : 수정된 것 표시 (상태보관)  
↓  
읽어왔는지 판별 중요

## □ Reading Uncommitted Data ("dirty reads")

failure

read한 게 dirty해졌다고 간주  
cf) write했지 않지만 실제 내용은 없는 경우

```
T1: BEGIN R(X), X=X-N, W(X), R(Y), Y=Y+N, W(Y) END
T2: BEGIN R(X), X=X+M, W(X) END
```

time →

- Let's say that  $X = 100$ ,  $Y = 0$  at the start and  $N = 50$ ,  $M = 10$ . Then after processing T1 and T2,
  - X should be 110.
  - Y should be 0.

# Anomalies with Interleaved Execution (Cont'd)

- Reading Uncommitted Data (“dirty reads”)  
Cont'd

T1: R(X), X=X-N, W(X)

T2:

R(X), X=X+M, W(X)

R(Y), fails then abort.

아직 undo

ex) 일관성 수준에 따라

time

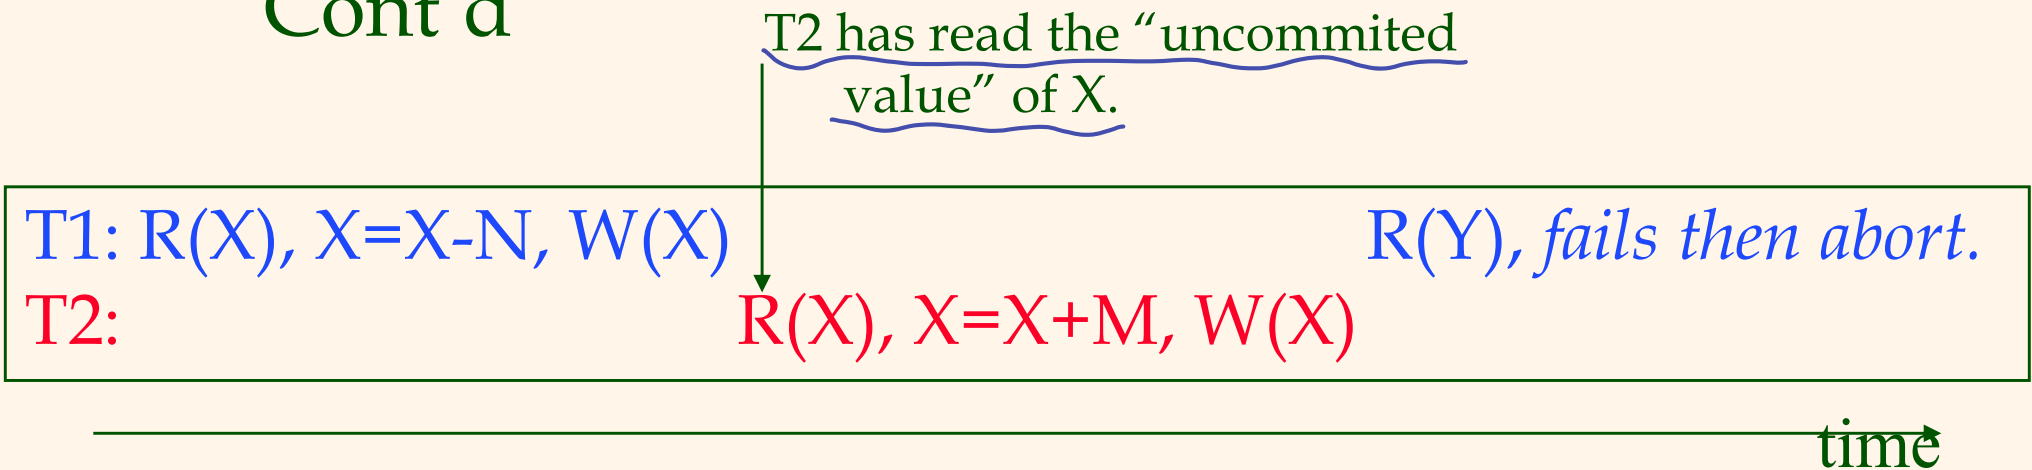
Let's say that  $X = 100$ ,  $Y=0$  at the start and  $N = 50$ ,  $M = 10$ .

X : 50 60

Y: 0

# Anomalies with Interleaved Execution (Cont'd)

- Reading Uncommitted Data (“dirty reads”)  
Cont'd



Let's say that  $X = 100$ ,  $Y=0$  at the start and  $N = 50$ ,  $M = 10$ .

X :	50	60	
Y:			0

↑  
 失敗했는데도 값이 바뀌었음  
 100이 아니라 60

# Anomalies with Interleaved Execution (Cont'd)

## ❑ Incorrect summary 결과값은 맞음

```
T1: BEGIN Sum=0, R(X), Sum=Sum+X, R(Y), Sum+=Y, print(Sum) END  
T2: BEGIN R(X), X=X-N, W(X), R(Y), Y=Y+N, W(Y) END
```

- ❑ T1 is a transaction that reports the sum of the balance of two accounts (X and Y).
- ❑ T2 is a transaction that transfers the balance (N) from an account (X) to the other account (Y)
- ❑ Let's say that  $X = 100$ ,  $Y = 0$  at the start and  $N = 50$ . Then after processing T1 and T2,
  - Sum should be either 100.

# Anomalies with Interleaved Execution (Cont'd)

## □ Incorrect summary (Cont'd)

T1:	Sum=0, R(X), Sum=Sum+X, R(Y), Sum+=Y, print(Sum)
T2:	R(X), X=X-N, W(X), R(Y), Y=Y+N, W(Y)

Let's say that  $X = 100$ ,  $Y=0$  at the start and  $N = 50$ .

X :	50				
Y:					50
Sum:	0	50	50	50	

# Anomalies with Interleaved Execution (Cont'd)

## □ Incorrect summary (Cont'd)

T1 reads X after N is  
subtracted but before N is  
added to Y

T1: Sum=0, R(X), Sum=Sum+X, R(Y), Sum+=Y, print(Sum)

T2: R(X), X=X-N, W(X),


R(Y), Y=Y+N, W(Y)

Let's say that  $X = 100$ ,  $Y=0$  at the start and  $N = 50$ .

X: 50

Y: 50

Sum: 0 50 50 50



# *Wait for a second .....*

- Why do we learn about several different kinds of anomalies ?
  - DBMS should allow multiple transactions can be processed concurrently.
  - Operations in multiple transactions can be interleaved.
  - Interleaving the operations may cause such anomalies.
  - So we need to have a concurrency control to avoid such anomalies.



# Atomicity of Transactions

- A transaction might *commit* after completing all its actions, or it could *abort* (or be aborted by the DBMS) after executing some actions.
- A very important property guaranteed by the DBMS for all transactions is that they are *atomic*. That is, a user can think of a transaction as
  1. always executing all its operations, or
  2. not executing any operations at all.
- What if something wrong during the processing of a transaction? In other words, what if some failure occurs in the middle of execution transactions.

) all or none

( $\frac{1}{2} \rightarrow \text{commit}$   
 $\frac{1}{2} \rightarrow \text{abort}$ )

11/21

# Recovery is needed in DBMS

dirty bit  
abort : 중단되면 세움으로


- Failures can happen in many forms:
  - Computer failure (system crash) : e.g. *hardware, network failure*
  - Transaction or system error : e.g. *divide by zero*
  - Local errors or exception conditions raised : e.g. *low balance*
  - Concurrency control enforcement : e.g. *not serializable*
  - Disk failure
  - Physical problems
- We need a **recovery** system in DBMS : DBMS **logs** all **actions** so that it *may undo* the actions of non-committed transactions.

# Transaction States and Transition Diagram

## □ Transaction States

- **Active:** a transaction is begun. The transaction can issue *Read* and *Write* operations.
- **Committed:** a transaction issued *Commit* operation.
- **Aborted:** a transaction failed or issued *Abort* operation. *un-do, rollback?*
- **Terminated:** a transaction is done and leaves the system.

## □ Transaction States Transition Diagram



# ACID Properties of Transaction

- **Atomicity** : a transaction is an atomic unit of processing. all or nothing
- **Consistency** : a transaction preserves the consistency of the database. → transition (변태) +) 수행 중에는? → 논외.
- **Isolation** : the execution of a transaction should not be interfered by any other transaction. cf) Isolation level
- **Durability** : changes made by any committed transaction must be persistent in the database.

# Schedule of Transactions and Serializability

## □ Explanation by example.

T1:	BEGIN	$A = A + 100$ ,	$B = B - 100$	END
T2:	BEGIN	$A = 1.05 * A$ ,	$B = 1.05 * B$	END

- Intuitively, the first transaction is transferring \$100 from B's account to A's account. The second is crediting both accounts with a 5% interest payment.
- There is no guarantee that T1 will execute before T2 or vice-versa, if both are submitted simultaneously. However, the net effect *had better* be equivalent to these two transactions running serially in some order.



# *Schedule of Transactions and Serializability (Cont'd)*

- What can you expect, assuming that you do not have any idea of interleaved executions and that  $A = B = 100$  before you execute those transactions ?

T1:	BEGIN	$A=A+100,$	$B=B-100$	END
T2:	BEGIN	$A=1.05*A,$	$B=1.05*B$	END

# Schedule of Transactions and Serializable Schedule (Cont'd)

□ You would expect:

T1:	BEGIN	A=A+100,	B=B-100	END
T2:	BEGIN	A=1.05*A,	B=1.05*B	END

□ If T1 is processed and then T2 is processed, then  $A = 210$ ,  $B = 0$ .

Serial schedule ①

T1:	BEGIN	A=A+100,	B=B-100	END		
T2:			BEGIN	A=1.05*A,	B=1.05*B	END

□ If T2 is processed and then T1 is processed, then  $A = 205$ ,  $B = 5$ .

T1:	Serial schedule ②		BEGIN	A=A+100,	B=B-100	END
T2:	BEGIN	A=1.05*A,	B=1.05*B	END		

# Schedule of Transactions and Serializable Schedule (Cont'd)

- Consider a possible interleaving (schedule A):

serializable(s)

T1:	$A = A + 100,$	$B = B - 100$
T2:	$A = 1.05 * A,$	$B = 1.05 * B$

- This is OK. Why ?

- But what about the following schedule B ?

serializable(X) : 위의 두 serial schedule 중 어떤 것과도 결과가 같지 않음

T1:	$A = A + 100,$	$B = B - 100$
T2:	$A = 1.05 * A, B = 1.05 * B$	





# Schedule of Transactions and Serializable Schedule (Cont'd)

serial vs. serializable

↳ 하나가 끝나야 다른 것 시작  
Interleave으로 실행한 결과가  
serial schedule과 같으면  
serializable schedule

- Serial schedule: Schedule that does not interleave the actions of different transactions.  
serial ↔ concurrent
  - T1 and T2
  - T2 and T1
- Serializable schedule: A schedule that is equivalent to some serial execution of the transactions.
  - In the previous slides, Schedule A is serializable schedule while Schedule B is not serializable schedule.



# Conflict Serializable Schedules

Two operations in a schedule are said to be **conflict** if:

- they belong to different transactions
- they access the same item
- at least one of the operations is **W(X)**.

→ 같은 transaction에 속해 있으면 conflict X

Schedule1:		
T1:	<u>R(X)</u> ,	R(Y) .....commit
T2:	<u>W(X)</u>	commit

Schedule1 (by conventional notation)		
R1(X),	↔	W2(X)
		R1(Y), <u>C2, C1</u>

동시 실행 가능  
가장 먼저 commit

commit

# Conflict Serializable Schedules (Cont'd)

$R_1(X) \ R_1(Y) \ W_2(X)$   
 $R_1(X) \ W_2(X) \ R_1(Y)$  } → conflict equivalent

conflict 측면에서 equivalent

□ Two schedules are **conflict equivalent** if:

- Involve the same operations of the same transactions → 순서는 달라도 내용은 같아야 함
- Every pair of conflicting operations is ordered the same way → conflict operation의 순서가 같아야 함

□ Schedule S is **conflict serializable** if S is conflict equivalent to some serial schedule.

Serial schedule과  
conflict equivalent 하다

# Example

< serial execution >


Serial 하게 실행하는 법 : 2가지

$\frac{T1 \rightarrow T2}{①}$  or  $\frac{T2 \rightarrow T1}{②}$

Serial schedule과  
Conflict 측면에서 같아야 함

≠ parallel  
interleave (동시 실행)

□ A schedule that is not conflict serializable:



즉이면 ①과 같아짐!  
→ conflict serializable

①  
 $T1 \rightarrow T2$

T1: R(A), W(A), R(B), W(B)

T2:

R(A), W(A), R(B), W(B)

②  
 $T2 \rightarrow T1$

T1:

R(A), W(A), R(B), W(B)

T2:

R(A), W(A), R(B), W(B)

⇒ 순서가 같은 게 많음!!  
Conflict serializable이 될  
수 없음

(백만개씩, 파랑끼리)

3쌍의 operation이 A에 대한 conflict  
두면  $T1 \rightarrow T2$

(백만개씩, 파랑끼리)  
3쌍의 operation이 B에 대한 conflict  
두면  $T2 \rightarrow T1$


transaction  $\rightarrow$  node  
conflict  $\rightarrow$  edge  
operation  $\rightarrow$  dependency graph

## Example (Cont'd)

- A schedule that is not conflict serializable:

Node  
2:11

T1:	R(A), W(A),	R(B), W(B)
T2:	R(A), W(A), R(B), W(B)	



"cycle이 있으면 conflict serializable이아님!"

Dependency graph

- The cycle in the graph reveals the problem. The output of T1 depends on T2, and vice-versa.

# Dependency Graph


transaction 개수가 많으면  
conflict serializable 한지 찾기 힘들다

- Dependency graph: One node per transaction;  
edge from  $T_i$  to  $T_j$  if an operation of  $T_i$   
precedes and conflicts with any operations of  
 $T_j$ .
- Theorem: Schedule is conflict serializable if  
and only if its dependency graph is acyclic.

cycle이 없다

필요충분조건  
(동치)

conflict serializable 한지 유지하는 법



서로 필요한 resource  
이때따른 매커니즘도 필요

access control mechanism

acquire = get

OS method (scheduling)  
누가 먼저 할 것인가  
(policy)  
FIFO? → queue

# Lock-Based Concurrency Control

거기서 거지자

cf) protocol : 계약, 협약  
방법

☆  
read lock = shared lock  
write lock = exclusive lock  
못간다면 기다려라

## Two-phase Locking (2PL) Protocol:

① Each transaction must obtain a **S (shared)** lock on object before reading, and an **X (exclusive)** lock on object before writing.   
↳ lock을 얻은 때까지 기다려라

② A transaction can not request additional locks once it releases any lock. (growing phase and shrinking phase)

③ If a transaction holds an X lock on an object, no other transaction can get a lock (S or X) on that object. Or no transaction can not get X lock on an object if other transaction holds a lock (S or X) on that object.


하나의 lock을  
가지고있으면  
동시성까지도

SS → 0  
SX  
XS  
XX  
→ X  
호환안됨

ex) T1이 A에 대한 X를 가지고있으면 T2는 A에 대한 어떤 lock도 가질 수 없음

\* R, W 둘다 하지 않을 것이라고  
DBMS가 아는 시점  
: commit 또는 rollback  
(transaction이 끝날때)

2PL allows only conflict serializable schedules.



# Lock-Based Concurrency Control (Con'd)

transaction 끝날때  
끝났다고 꼭 꼭 말하셈!

→ Strict 2PL  
DBM이 마지막 오퍼레이션이  
원치 못하니까 일반 2PLX

→ 경쟁하고 놓지 않으면  
concurrency level ↓  
→ 2PL에 비해 Strict 2PL이 낫음

## □ Strict Two-phase Locking (Strict 2PL) Protocol:

- Each transaction must obtain a S (*shared*) lock on object before reading, and an X (*exclusive*) lock on object before writing.
- All locks held by a transaction are released when the transaction completes.
- If a transaction holds an X lock on an object, no other transaction can get a lock (S or X) on that object. Or no transaction can not get X lock on an object if other transaction holds a lock (S or X) on that object.

□ Strict 2PL allows only conflict serializable schedules.



# Aborting a Transaction

- If a transaction  $T_i$  is aborted, all its actions have to be undone. Not only that, if  $T_j$  reads an object last written by  $T_i$ ,  $T_j$  must be aborted as well!
- Most systems avoid such <sup>상호</sup>cascading aborts by releasing a transaction's locks only at commit time.
  - If  $T_i$  writes an object,  $T_j$  can read this only after  $T_i$  commits.

T1: R(X), X=X-N, W(X)

..abort.

T2:

R(X), X=X+M, this must be aborted too.

↓  
T1 작업 후에 이 데이터가 읽혀야 함 (종속)  
⇒ T1 abort 시 같이 undo

cf) deadlock detection / recovery

## Aborting a Transaction (Cont'd)

- In order to *undo* the actions of an aborted transaction, the DBMS maintains a *log* in which every write is recorded. This mechanism is also used to recover from system crashes: all active transactions at the time of the crash are aborted when the system comes back up.

T1: R(X), X=X-N, W(X)

.... fails then abort.

Log: T1 updates  
X from 100 to 50.

# Summary

- Concurrency control and recovery are among the most important functions provided by a DBMS.
- Users need not worry about concurrency!!!!!!
  - DBMS does it for you !!!!!
  - For example, DBMS automatically inserts lock/unlock requests and schedules actions of different transactions in such a way as to ensure that the resulting execution is equivalent to executing the transactions one after the other in some order. + recovery, system reboot, ... 자동화
- Recovery is used to undo the actions of aborted transactions and to restore the system to a consistent state after a crash.
- Users need not worry about recovery!!!!!!
  - DBMS does it for you !!!!!
  - For example, DBMS maintains the log.

마지막 수업 날 강의 요약

11/30