


교수님명충원안되면...

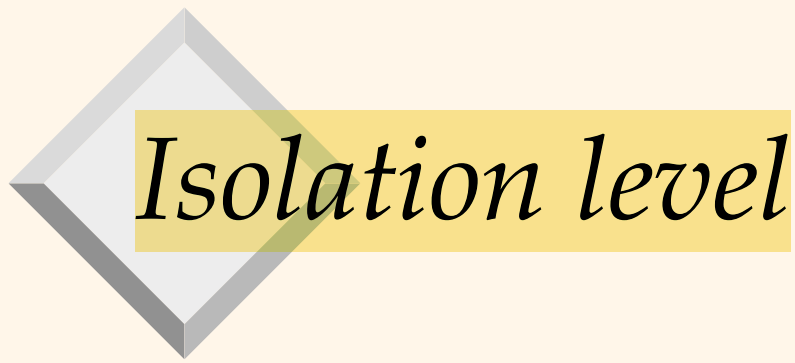


# *Crash Recovery*



# Review: *The ACID properties*

- ❑ **A**tomicity: All actions in the transaction happen, or none happen.
- ❑ **C**onsistency: If each transaction is consistent, and the DB starts consistent, it ends up consistent.
- ❑ **I**solation: Execution of one transaction is isolated from that of other transactions.
- ❑ **D**urability: If a transaction commits, its effects persist.
- ❑ The **Recovery Manager** guarantees Atomicity & Durability.



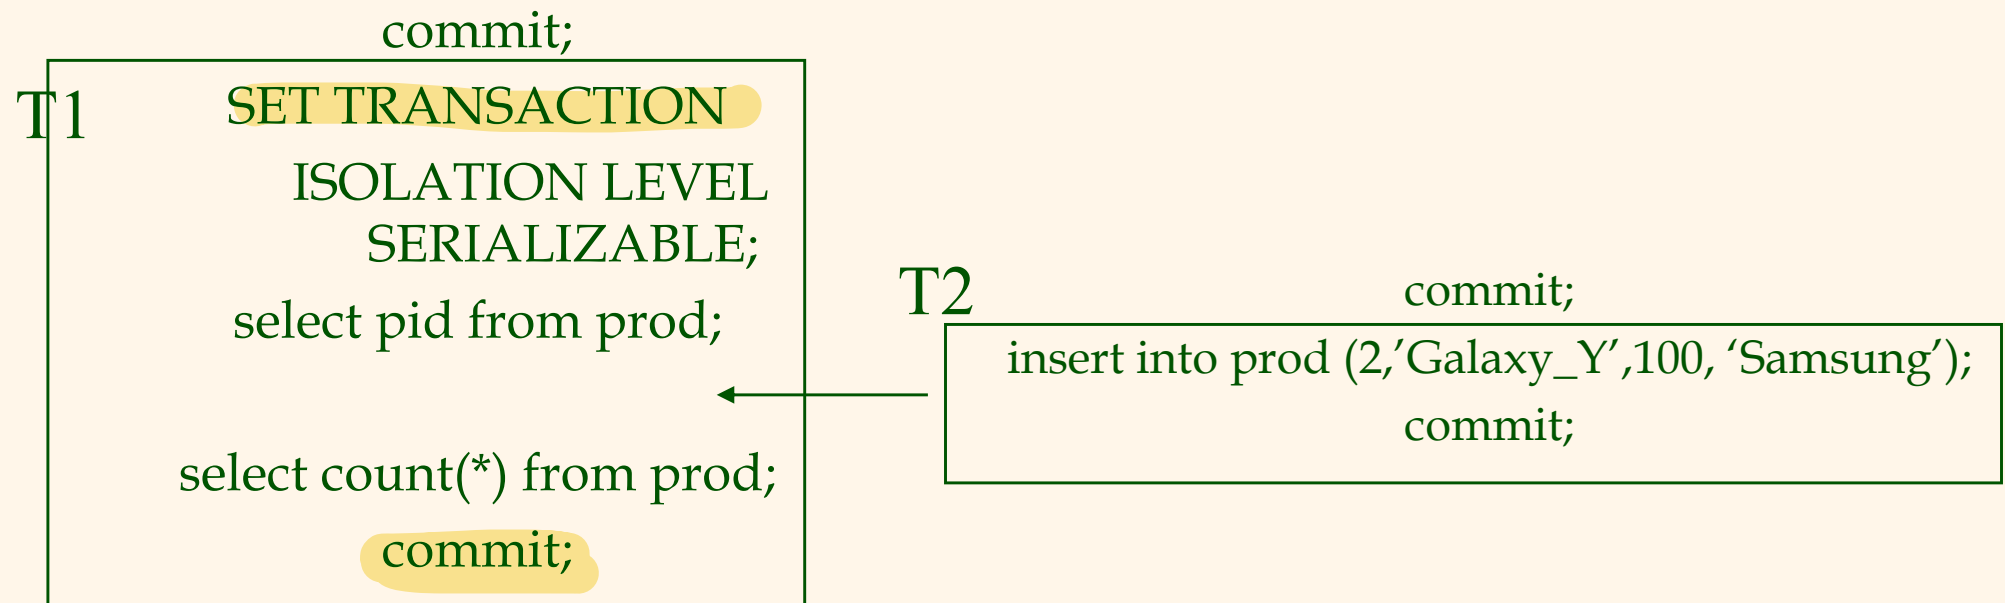
# *Isolation level*

- Isolation level of a transaction
  - Can be set in `SET TRANSACTION` statement
  - *Serializable*: default in SQL standard.
  - *Repeatable Read*: prevents non-repeatable read.
  - *Read Committed*: default in Oracle DBMS.
    - See changes only committed by another transactions.
    - Prevents dirty-read anomaly.
  - *Read Uncommitted*:
    - See changes incurred by any (including uncommitted) transactions.



# Isolation level Cont'd

## □ Serializable

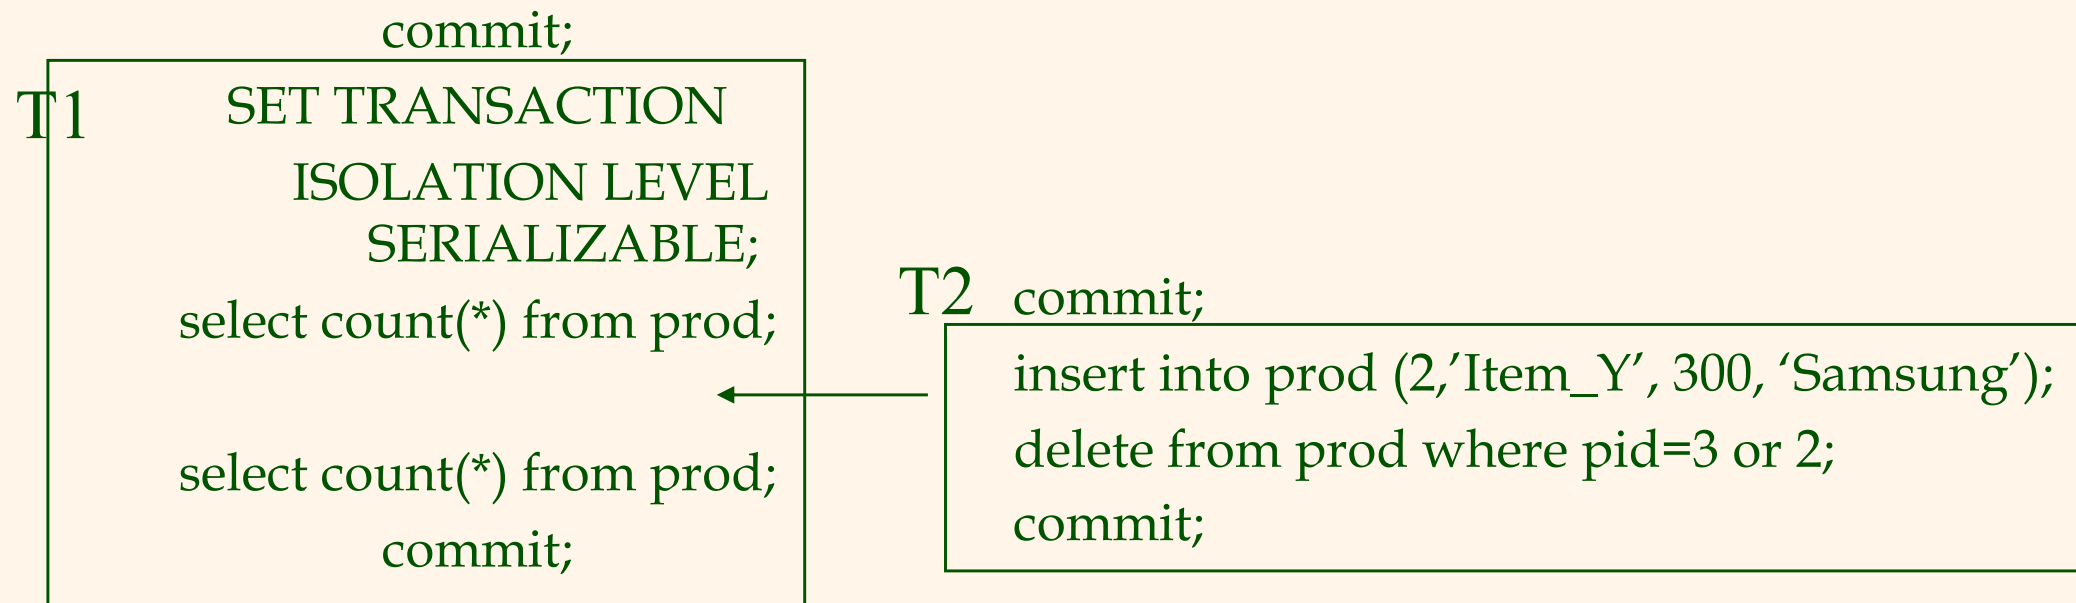


- If DBMS supports the serializability, then the result should be the same as either T1 and T2, or T2 and T1.
- Lock the proj table in S mode so that any write operation to the proj table is not allowed. => prevent *phantom*.




# Isolation level Cont'd

## □ Serializable cont'd




- If DBMS supports the serializability, then the result should be the same as either T1 and T2, or T2 and T1.
- Lock the proj table in **S mode** so that any write operation to the proj table is not allowed. => prevent **phantom**.




# Isolation level Cont'd

## □ Repeatable Read

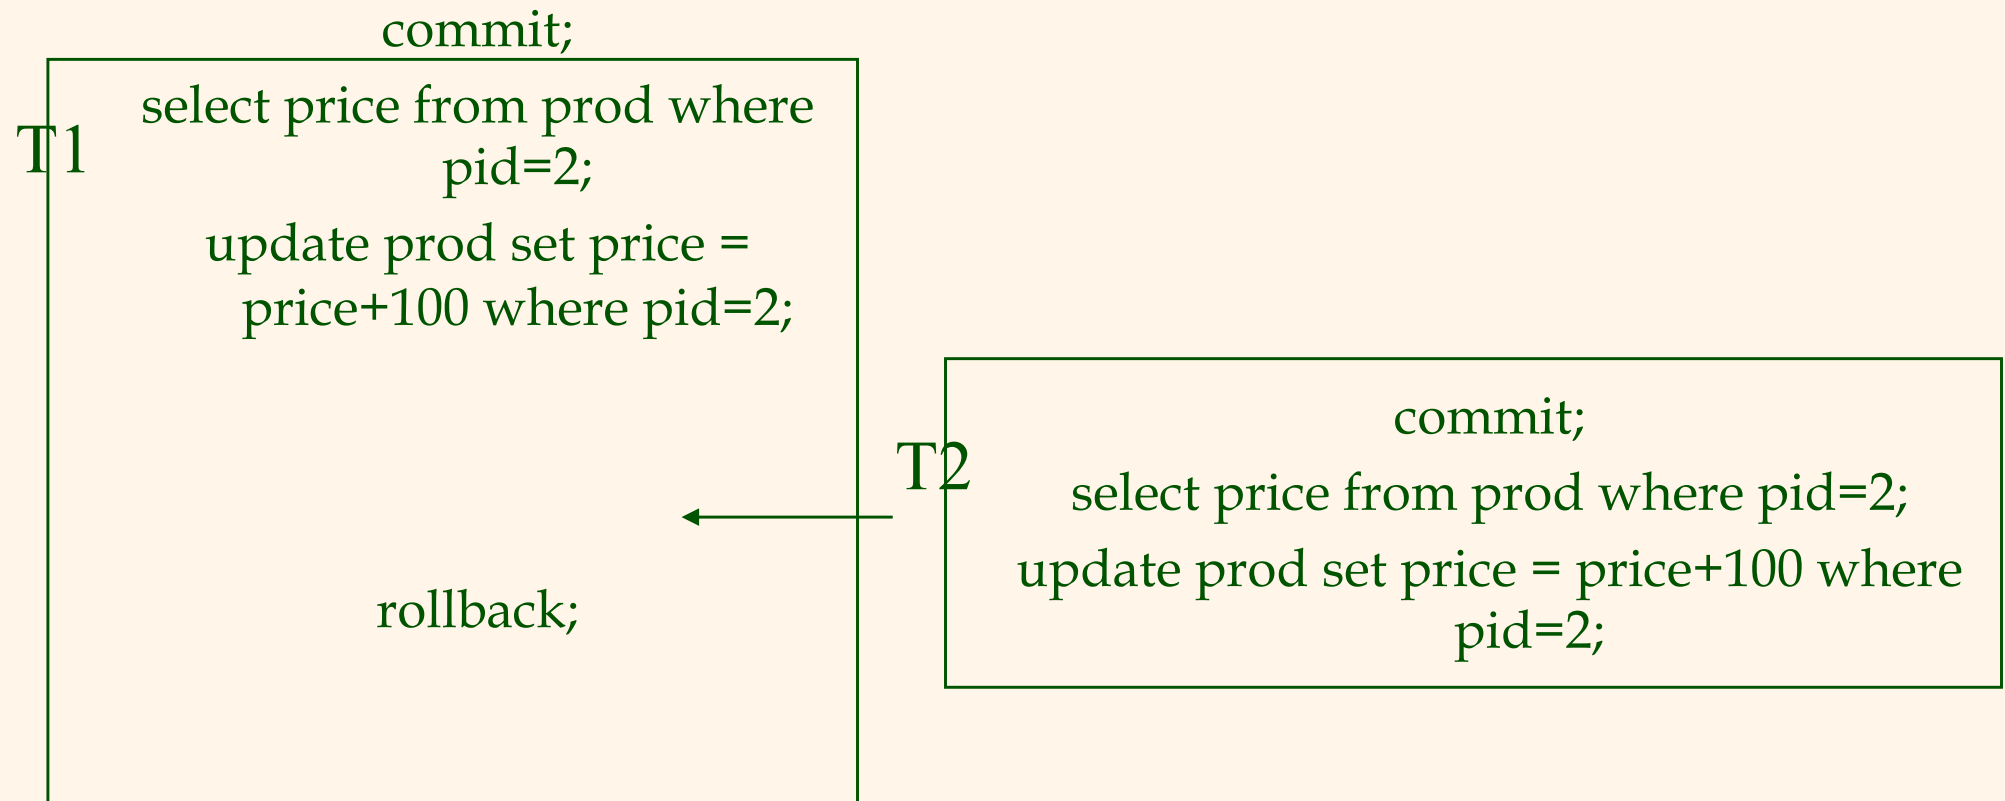


- Non repeatable read : the budget value of the first select is different to the budget value of the second select.
- If the isolation level is set to REPEATABLE READ, then two budget values are the same. That's why it is called *repeatable read*.



# Isolation level Cont'd

## □ Read Committed




- **Dirty read**: T2 read the budget written by T1 which has not committed. => Not desired, No practical at all.

# Motivation

- Atomicity:
  - Transactions may abort (“Rollback”).
- Durability:
  - What if DBMS stops running? (Causes?)

- Desired Behavior after system restarts:
  - T1, T2 & T3 should be durable.
  - T4 & T5 should be aborted (effects not seen).





# Recovery Strategy

디스크 업데이트는 커밋되면

## □ Deferred update

- Do not physically update the database in disk until the transaction commit.
- During the commit, the updates are first recorded persistently in the log and then written to the database.
- No-UNDO/REDO

DB 업데이트는 W하면

## □ Immediate update

- the database may be “updated” immediately after the write operations although the transaction does not reach commit.
- What’s happening in “updating”? See the following slide.

# Handling the Buffer Pool

→ force stealing 정책!

- **Force** committed all updates to disk when the transaction commits.


- Poor response time.
- Providing durability is not hard.

- **Steal** buffer frames from uncommitted transactions.

- If not, poor throughput.
- Providing atomicity is not easy.


	No Steal	Steal
Force	Trivial	
No Force		Desired/ in Practice

# *Basic Idea: Logging*



- **REDO** : It may be necessary to “redo” the operations of a committed transaction
  - Need new values.
- **UNDO** : We need “undo” the operations of a failed transaction or of another uncommitted transaction which reads the dirty data.
  - Need old values.

# Basic Idea: Logging




- Record REDO and UNDO information, for every update, in a *log*.
  - Sequential writes to log (put it on a separate disk).
  - Minimal info (diff) written to log, so multiple updates fit in a single log page.
- Log: An ordered list of REDO/UNDO actions
  - Log record contains:
    - <transactionID, data\_item, old value, new value>
    - and other info such as begin, and commit/rollback.
  - Can vary according to the recovery scheme.

# Write-Ahead Logging (WAL)

- The Write-Ahead Logging Protocol:
  - ① Must **force** the **log record** for an update before the corresponding **data page** gets to disk.
  - ② Must **write all log records** for a transaction before commit.
- #1 guarantees **Atomicity**.
- #2 guarantees **Durability**.
- Exactly how is logging (and recovery!) done?
  - Like **ARIES**(Algorithm for Recovery and Isolation Exploiting Semantics) by IBM Almaden Research.
  - IBM DB2, Informix, MS SQL Server, Oracle 8, Sybase uses ARIES or its variant.

# Big Picture



다음페이지!

- T3 and T5 must be undone. Why? → commit 되지 않음  
→ system failure (atomicity를위해)
- T2 and T4 must be redone. Why? → checkpoint 이후에 commit, 제대로 했다는보장이없어서 redo함
- How about T1? → redo 필요x, checkpoint 이전에 commit



# Big Picture

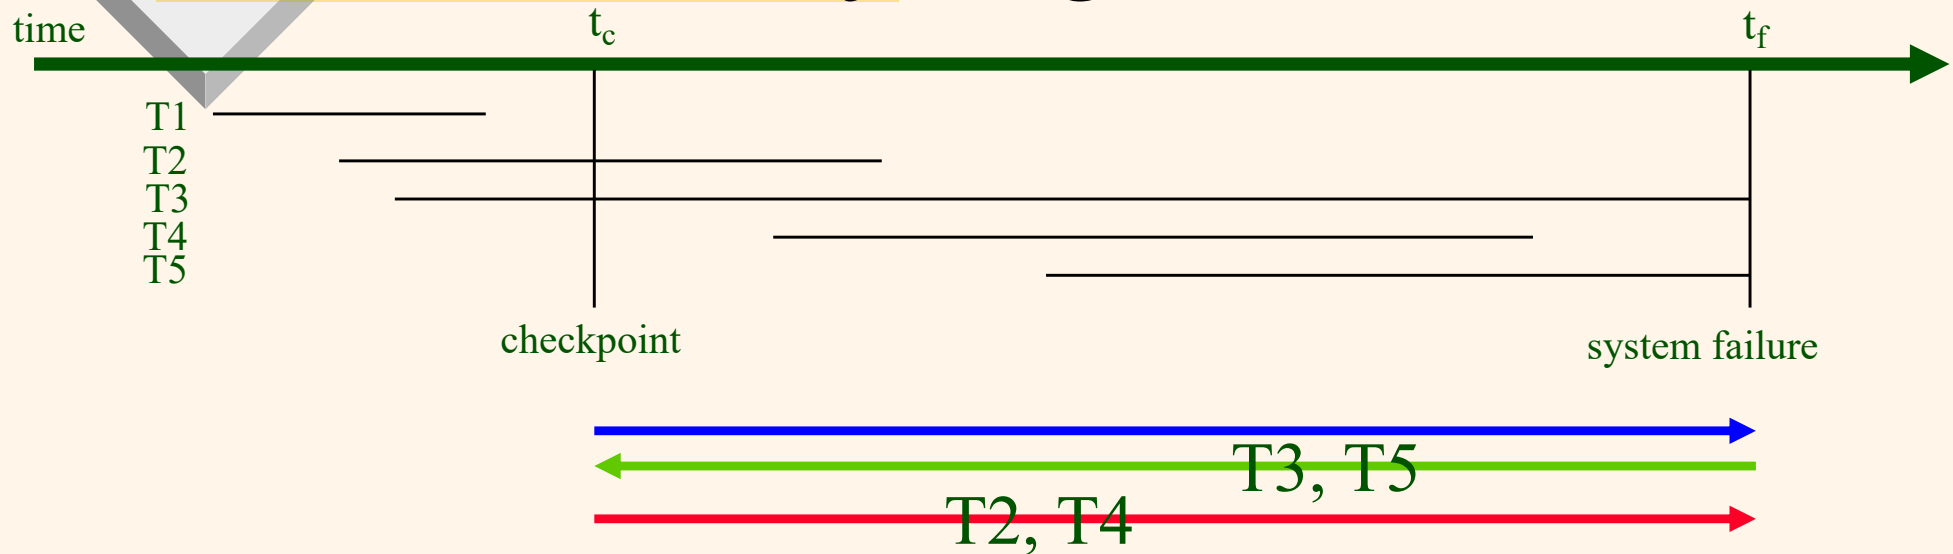
- T3 and T5 must be undone. Why?
  - Any change that was made by **uncommitted transaction** must be undone.
- T2 and T4 must be redone. Why?
  - No-force. In other words, there is no guarantee that their updates were actually written to the database.
- How about T1?
  - Updates were forced out to the database at  $t_c$  when the checkpoint was taken.

# Checkpointing

- Periodically, the DBMS creates a checkpoint, in order to minimize the time taken to recover in the event of a system crash. The checkpoint may involve in general :
  - Forcing a “checkpoint record” out to the log storage.
  - Forcing the content of database buffers out to the database.
  - Writing the the address of the checkpoint record within the log into a “master record”. ex) OS에서 첫 부팅때 boot record
- No work done prior to the checkpoint need ever be redone.




# Crash Recovery: Big Picture



- Start from a **checkpoint** (found via **master** record).
  - Three phases. Need to:
    - ① Figure out which Xacts committed since checkpoint, which failed (**Analysis**). 순방향으로 가면서 undo/redo 대상 분류
    - ② **UNDO** effects of failed Xacts 역방향으로 가면서 undo
    - ③ **REDO** all actions 순방향으로 가면서 redo
- ※ 반드시 undo 먼저 하지 않아도 결과는 같음  
 sequential access :  $\rightarrow$  효율적  
 random access가 된다면 시간 낭비가 많음

# Crash Recovery: Big Picture

↗ recursive하게 여러 번 해도  
같은 결과를 내야 함.




UNDO-list : {T2,T3}	{T2,T3,T4}	{T3,T4,T5}	{T3,T5}
REDO-list : {}	{T2}		{T2,T4}

## □ Analysis Phase

- Initializes UNDO-list to have all transactions listed in the checkpoint record. Initializes REDO-list empty.
- If finding a BEGIN TRANSACTION record for a transaction, it adds that transaction to the UNDO-list.
- If finding a COMMIT TRANSACTION record for a transaction, it moves that transaction from UNDO-list to the REDO-list.

# Crash Recovery: Big Picture



UNDO-list : {T3, T5}


REDO-list : {T2, T4}

## □ Undo Phase

□ Undoing the transactions in the undo-list

## □ Redo Phase

□ Redoing the transactions in the redo-list



# *Exactly how is logging (and recovery!) done?*

- Like **ARIES**(Algorithm for Recovery and Isolation Exploiting Semantics) by IBM Almaden Research.
- IBM DB2, Informix, MS SQL Server, Oracle 8, Sybase uses ARIES or similar to its variant.
- Should I know ARIES Algorithm??????



# Summary of Logging/Recovery

기록에 제대로 남겨야!

- **Recovery Manager** guarantees Atomicity & Durability.
- Use WAL to allow STEAL/NO-FORCE w/o sacrificing correctness.
- **Checkpointing**: A quick way to limit the amount of log to scan on recovery.

- Recovery works in 3 phases:

- Analysis:
- Redo:
- Undo:

concurrency control & recovery

↓  
scheduling  
→ locking

↓  
atomicity & durability  
→ logging

기말 12/7