
Chapter 3

Finding Similar Items

Finding “Similar” Items

- Many problems can be expressed as finding *similar* items
 - Find nearest neighbors in *high dimensional* space
 - One of the fundamental data mining problems
- Examples
 - Finding near-duplicate Web pages
 - Plagiarisms or mirrors
 - Finding pages with similar words
 - Duplicate detection, classification by topic
 - Finding customers who purchased similar products
 - Products with similar customers (recommender systems)
 - Finding images with similar features
 - Similarly, users who visited similar websites



Problem for This Chapter

■ Given

- High dimensional data points x_1, x_2, \dots
 - (ex) Image: a long vector of pixel colors
- Some distance function $d(x_1, x_2)$
 - Which quantifies the “distance” between x_1 and x_2

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 2 & 1 \\ 0 & 1 & 0 \end{bmatrix} \rightarrow [1 \ 2 \ 1 \ 0 \ 2 \ 1 \ 0 \ 1 \ 0]$$

■ Goal

- Find all pairs of data points (x_i, x_j) that are within some distance threshold $d(x_i, x_j) \leq s$

■ Note

- Naïve solution would take $O(N^2)$ where N is the number of data points
- How can this be done in $O(N)$??

Finding “Similar” Documents

- Here, we focus on finding *similar documents*
- Goal
 - Given a **large** (e.g., 10^9) number of documents, find “near duplicate” pairs
- Applications
 - Mirror websites (don’t want to show both in search results)
 - Similar news articles (cluster articles by “same story”)
- Difficulties
 - Many small pieces of one document can appear out of order in another
 - There are too many documents to compare all pairs
 - Documents are so large or so many that they cannot fit in main memory

3 Steps for Finding Similar Documents

1. Shingling

- Convert documents to ***sets***

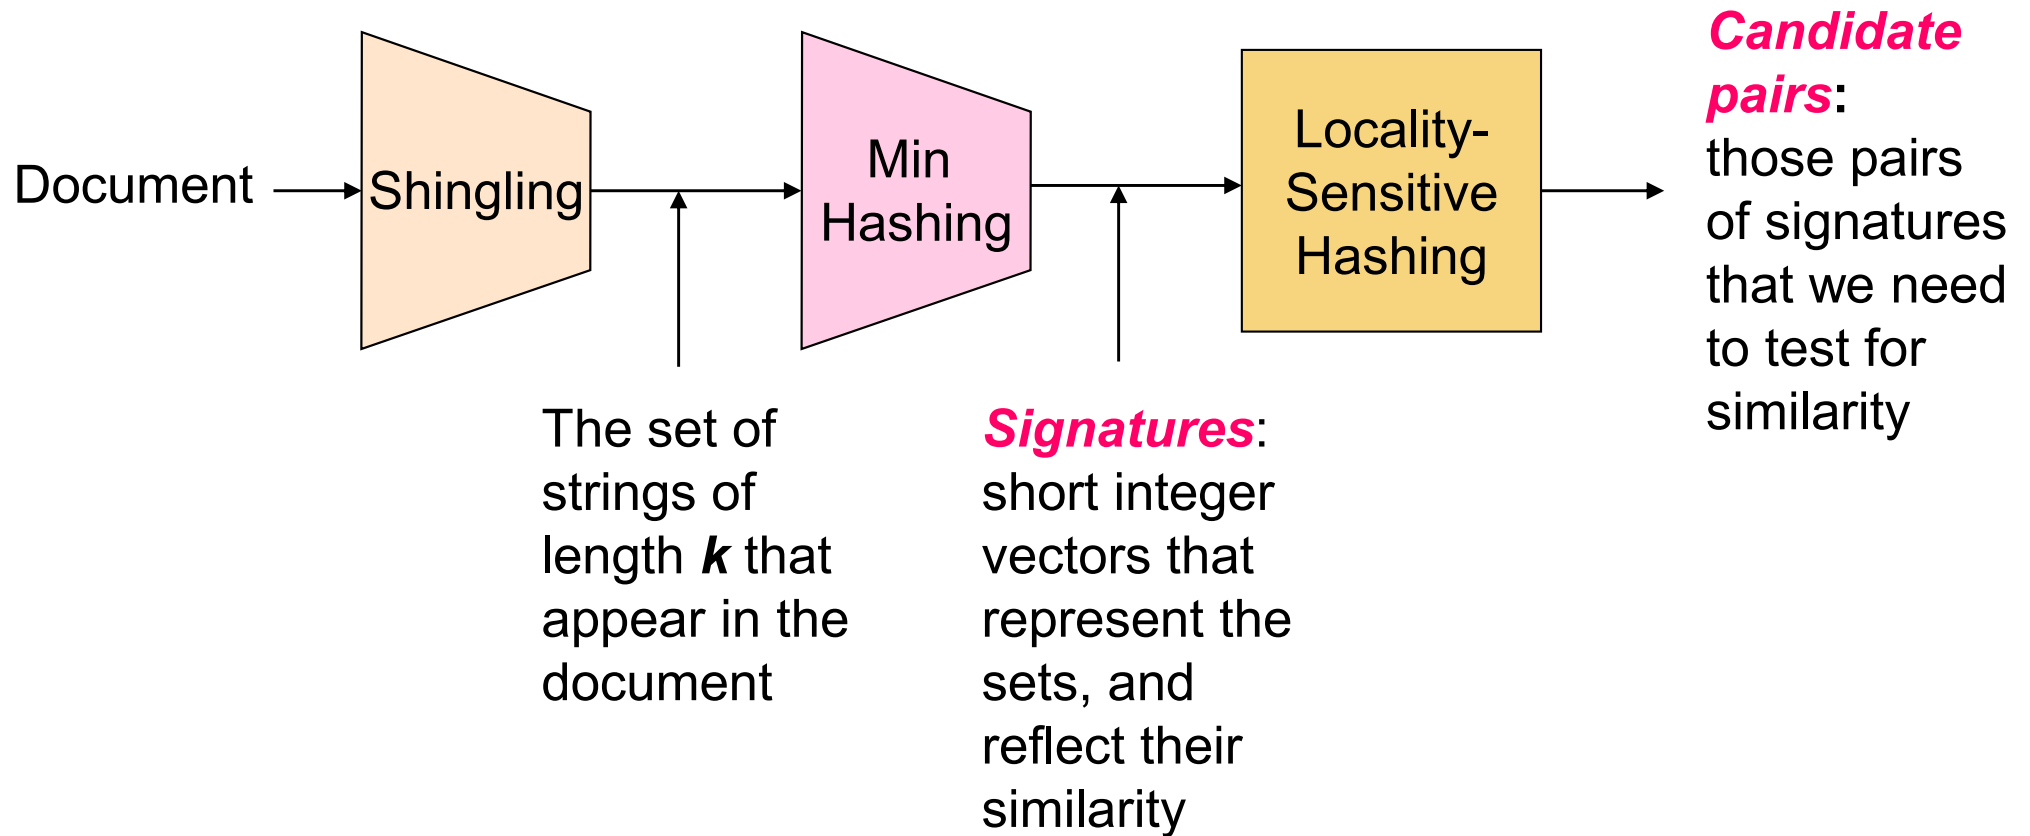
2. Minhashing

- Convert large sets to short ***signatures***, while preserving similarity

3. Locality-Sensitive Hashing

- Focus on pairs of signatures likely to be from similar documents
- The results are ***candidate pairs***

The Big Picture




Jaccard Similarity

- The similarity of **sets** by looking at the size of their intersection
- The Jaccard similarity of sets S and T

$$SIM(S, T) = \frac{|S \cap T|}{|S \cup T|}$$

- Example

- $SIM(S, T) = 3/8$



Similarity of Documents

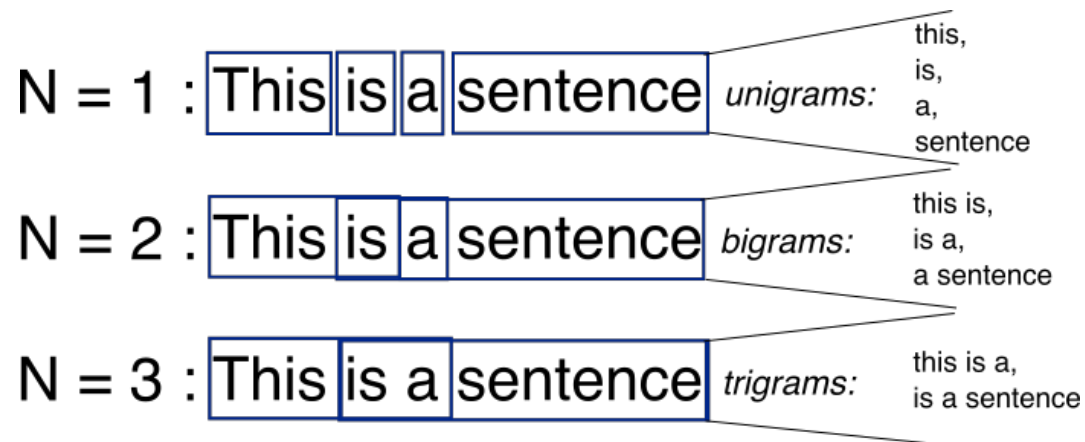
- Here, we focus on *character-level (i.e., textual)* similarity
 - Note that “similar meaning” requires other techniques
- Testing whether two documents are exact duplicates is easy
 - However, in many applications, the documents are **not** identical
- Applications of textual similarity
 - Plagiarism
 - Mirror pages
 - Search engines should avoid showing two pages that nearly identical
 - Articles from the same source
 - News aggregators(e.g., Google News) should show only one for each article

Collaborative Filtering

- A process whereby we recommend to users items that were liked by other users who have exhibited similar tastes
 - Another class of applications where ***similarity of sets*** is very important
- Examples
 - On-line purchases (e.g., Amazon.com)
 - Two customers are similar if their sets of purchased items have a high Jaccard similarity
 - Movie ratings (e.g., NetFlix)
 - Moves are similar if they were rented or rated highly by many of the same customers
 - Customers are similar if they rented or rated highly many of the same movies

Documents as Sets

- Simple approaches
 - Document = set of words appearing in document
 - Document = set of “important” words
 - Don’t work well for this application. **Why?**
- Need to account for **ordering** of words!
- A different way: Shingles (or grams)!



Shingling of Documents

- The most effective way to represent documents as *sets*
 - For the purpose of identifying lexically similar documents
- Construct from the documents *the set of short strings* that appear within it
 - (ex) Document = `abccab` → the set of 3-shingles = $\{abc, bca, cab\}$



- Documents that share sentences (or phrases) will have many *common* elements in their shingling sets
 - Even if they appear in different orders in the two documents

Definition: k -Shingles (or n -Grams)

- A k -shingle for a document
 - Any substring of length k found within the document
- Then, we associate with each document the set of k -shingles that appear one or more times within that document
- Example
 - Document $D = \text{abcdabd}$
 - The set of 2-shingles for $D = \{\text{ab}, \text{bc}, \text{cd}, \text{cd}, \text{da}, \text{bd}\}$
 - A variation of shingling produces a bag, rather than a set

Treating White Space

- White space

- Blank, tab, newline, etc.

- Options

- Replace any sequence of white-space characters by a single blank
- We may *eliminate* whitespace altogether

- Example

- D_1 = “The plane was ready for touch down”
- D_2 = “The quarterback scored a touchdown”
- If we use $k = 9$ and retain the blanks, D_1 has shingles touch dow and ouch down, while D_2 has touchdown
- If we eliminated the blanks, both would have touchdown

Choosing the Shingle Size

- If we pick k too small
 - Even the documents that have no the same sentences or even phrases would have **high** Jaccard similarity (e.g., $k = 1$)
- Rule of thumb
 - k should be pick **large** enough that the probability of any given shingle appearing in any given document is low
- Example
 - $k = 5$ is ok for short documents (e.g., emails)
 - $k = 10$ is better for long documents (e.g., research articles)

Compressing Shingles

- We can **compress** long shingles by using a **hash function** h
 - h maps each k -shingle to $0, \dots, B - 1$, where B is the number of buckets
 - Then, a document is represented as the set of **hash values** of its k -shingles
- Not only the data been compacted, but we can now manipulate (hashed) shingles by single-word machine operations
- Example
 - $k = 2$, document $D = \text{ab cab}$
 - The set of 2-shingles = $\{\text{ab}, \text{bc}, \text{ca}\}$ ($9 \times 3 = 27$ bytes)
 - The set of their hash values = $\{1, 5, 7\}$ ($4 \times 3 = 12$ bytes)

Shingles Built from Words

- In many applications, we want to ignore stop words
 - (ex) “a,” “and,” “for,” etc.
- However, for the problem of finding similar news articles, defining a shingle to be ***a stop word followed by the next two words*** forms a useful set of shingles
 - Bias the set of shingles in favor of the article, rather than its surrounding material
- Example
 - An ad: “Buy Sudzo”
 - A news article: “**A** spokesperson **for the** Sudzo Corporation revealed today **that** studies **have** shown **it is** good **for** people **to** buy Sudzo Product”
 - The set of shingles = {“A spokesperson for”, “for the Sudzo”,}
 - Note that **none** are from the ad

Motivation for Minhash/LSH

- Suppose we need to find near-duplicate documents among $N = 1$ million documents
- Naïvely, we would have to compute *pairwise* Jaccard similarities for every pair of docs
 - $N(N-1)/2 \approx 5 \times 10^{11}$ comparisons
 - At 10^5 secs/day and 10^6 comparisons/sec, it would take 5 days
- For $N = 10$ million, it takes more than a year...

Similarity-Preserving Summaries of Sets

- Even if we hash each shingle to 4 bytes, the space needed to store all sets of shingles is still large
 - We may have millions of documents
- Our goal
 - Replace large sets by much smaller representations called *signatures*
 - (ex) 200,000 byte hashed-shingle sets → 1,000 byte signatures
- Important property required for signatures
 - We must be able to estimate the Jaccard similarity of two sets from their signatures *alone*
 - Note that it is not possible that the signatures give the exact similarity

Matrix Representation of Sets

- Characteristic matrix

- Columns: sets (documents)
- Rows: elements (shingles)
- 1 in row e and column s , if and only if e is a member of s

- Example

- $S_1 = \{a, d\}$, $S_2 = \{c\}$, $S_3 = \{b, d, e\}$, $S_4 = \{a, c, d\}$

<i>Element</i>	S_1	S_2	S_3	S_4
a	1	0	0	1
b	0	0	1	0
c	0	1	0	1
d	1	0	1	1
e	0	0	1	0

Minhashing (1/2)

- **Goal:** Find a hash function h such that
 - If $SIM(S_1, S_2)$ is **high**, then $h(S_1) = h(S_2)$ with a **high** probability
 - If $SIM(S_1, S_2)$ is **low**, then $h(S_1) \neq h(S_2)$ with a **high** probability
- Clearly, the hash function depends on the similarity metric
 - Not all similarity metrics have a suitable hash function
- There is a suitable hash function for the Jaccard similarity
 - It is called **minhashing**

Minhashing (2/2)

- Pick a random permutation of the rows of the characteristic matrix
 - (ex) $abcde \rightarrow beadc$
- Minhash function $h(S)$ for a set S
 - The index of the first row, **in the permuted order**, in which the column has 1

- Example

- Permuted order = $beadc$
- $h(S_1) = a$
- $h(S_2) = c$
- $h(S_3) = b$
- $h(S_4) = a$

<i>Element</i>	S_1	S_2	S_3	S_4
<i>b</i>	0	0	1	0
<i>e</i>	0	0	1	0
<i>a</i>	1	0	0	1
<i>d</i>	1	0	1	1
<i>c</i>	0	1	0	1

Minhashing and Jaccard Similarity

- Let S_1 and S_2 be two sets
- Let $SIM(S_1, S_2)$ be the Jaccard similarity of S_1 and S_2
- Let h be the minhash function for a random permutation of rows
- ***Then, the probability that $h(S_1) = h(S_2)$ equals $SIM(S_1, S_2)$***

Proof (1/2)

- Consider the columns for sets S_1 and S_2

<i>Element</i>	S_1	S_2	
b	1	0	→ Type Y
e	1	0	
a	0	1	
d	1	1	→ Type X
c	0	1	

- Type X rows have 1 in both columns
 - Type Y rows have 1 in one of the columns and 0 in the other
 - Type Z rows have 0 in both columns
- Let x and y be the number of rows of type X and Y , respectively
- Then, $SIM(S_1, S_2) = x/(x + y)$

Proof (2/2)

- Now consider the probability that $h(S_1) = h(S_2)$
- Suppose we proceed from the top the rows permuted randomly
- The probability that we shall meet a type X row before we meet a type Y row
 - $x/(x + y)$
 - This corresponds to the probability that $h(S_1) = h(S_2)$
- Therefore, the probability that $h(S_1) = h(S_2)$ is $x/(x + y)$, which is also $SIM(S_1, S_2)$

Minhash Signatures

- Suppose we represent sets by their characteristic matrix M
- We pick at n random permutations of the rows of M
 - $n = 100$ or several hundreds
- Let h_1, h_2, \dots, h_n be the minhash functions for n permutations
- The minhash signature for S is the vector $[h_1(S), h_2(S), \dots, h_n(S)]$
- Thus, we can form a signature matrix from M
 - The i th column of M is replaced by the minhash signature for the i th column

(Ex) Minhash Signatures

Characteristic matrix M

3 Permutations

This row becomes the 3rd row in the 3rd permutation

s_1	s_2	s_3	s_4
1	0	1	0
1	0	0	1
0	1	0	1
0	1	0	1
0	1	0	1
1	0	1	0
1	0	1	0

2	4	3
3	2	4
7	1	7
6	3	2
1	6	6
5	7	1
4	5	5

This row becomes the 4th row in the 1st permutation

Signature matrix

s_1	s_2	s_3	s_4
2	1	2	1
2	1	4	1
1	2	1	2

Computing Minhash Signatures

- Picking a random permutation and permuting rows is ***not*** feasible
 - Permuting millions or billions of rows is time-consuming
- Instead, we can simulate a random permutation by a ***randomly chosen*** hash function
 - A hash function h that maps $0, \dots, k-1$ to bucket numbers $0, \dots, k-1$
 - (ex) $h(x) = (ax + b) \bmod k$, where a and b are random integers
 - h “permutes” row r to position $h(r)$ in the permuted order
- Thus, instead of picking n random permutations of rows, we pick n randomly chosen hash functions h_1, h_2, \dots, h_n on the rows
 - The signature matrix is then constructed by considering each row in their given order

Implementation

- $SIG(i, c)$
 - The element of the signature matrix for the i th permutation and column c

s_1	s_2	s_3	s_4
2	1	2	1
2	1	4	1
1	2	1	2

----- $SIG(2, 4)$

- Algorithm

Initially, set $SIG(i, c)$ to ∞ for all i and c

For each row r do

 Compute $h_1(r), h_2(r), \dots, h_n(r)$

 For each column c do

 (a) If c has 0 in row r , do nothing

 (b) If c has 1 in row r , then for each $i = 1, 2, \dots, n$
 set $SIG(i, c)$ to the smaller of the current value
 of $SIG(i, c)$ and $h_i(r)$

(Ex) Signature Matrix (1/7)


- Consider the following characteristic matrix and two hash functions: $h_1(x) = x + 1 \bmod 5$ and $h_2(x) = 3x + 1 \bmod 5$

Row	S_1	S_2	S_3	S_4	$x + 1 \bmod 5$	$3x + 1 \bmod 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

- Initially, we set $SIG(i, c)$ to ∞ for all i and c

	S_1	S_2	S_3	S_4
h_1	∞	∞	∞	∞
h_2	∞	∞	∞	∞

(Ex) Signature Matrix (2/7)

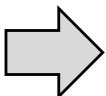


Row	S_1	S_2	S_3	S_4	$x + 1 \pmod 5$	$3x + 1 \pmod 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3


- First we consider row 0

- $h_1(0) = 1, h_2(0) = 1$

	S_1	S_2	S_3	S_4
h_1	∞	∞	∞	∞
h_2	∞	∞	∞	∞




	S_1	S_2	S_3	S_4
h_1	1	∞	∞	1
h_2	1	∞	∞	1



For S_1 , row **1** (originally row 0) is the first row whose column is 1 **for now**

(Ex) Signature Matrix (3/7)

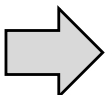


Row	S_1	S_2	S_3	S_4	$x + 1 \pmod{5}$	$3x + 1 \pmod{5}$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

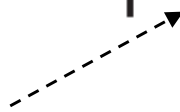
- Next we consider row 1

- $h_1(1) = 2, h_2(1) = 4$

	S_1	S_2	S_3	S_4
h_1	1	∞	∞	1
h_2	1	∞	∞	1




	S_1	S_2	S_3	S_4
h_1	1	∞	2	1
h_2	1	∞	4	1



For S_3 , row **4** (originally row 1) is the first row whose column is 1 **for now**

(Ex) Signature Matrix (4/7)

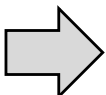


Row	S_1	S_2	S_3	S_4	$x + 1 \pmod 5$	$3x + 1 \pmod 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

- Next we consider row 2


- $h_1(2) = 3, h_2(2) = 2$

	S_1	S_2	S_3	S_4
h_1	1	∞	2	1
h_2	1	∞	4	1



	S_1	S_2	S_3	S_4
h_1	1	3	2	1
h_2	1	2	4	1

(Ex) Signature Matrix (5/7)

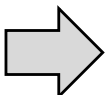


Row	S_1	S_2	S_3	S_4	$x + 1 \pmod{5}$	$3x + 1 \pmod{5}$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

- Next we consider row 3


- $h_1(3) = 4, h_2(3) = 0$

	S_1	S_2	S_3	S_4
h_1	1	3	2	1
h_2	1	2	4	1



	S_1	S_2	S_3	S_4
h_1	1	3	2	1
h_2	0	2	0	0

(Ex) Signature Matrix (6/7)

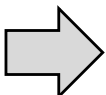


Row	S_1	S_2	S_3	S_4	$x + 1 \pmod{5}$	$3x + 1 \pmod{5}$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

- Next we consider row 4

– $h_1(3) = 0, h_2(3) = 3$

	S_1	S_2	S_3	S_4
h_1	1	3	2	1
h_2	0	2	0	0



	S_1	S_2	S_3	S_4
h_1	1	3	0	1
h_2	0	2	0	0

(Ex) Signature Matrix (7/7)

- Estimating Jaccard similarities from the signature matrix

Row	S_1	S_2	S_3	S_4
0	1	0	0	1
1	0	0	1	0
2	0	1	0	1
3	1	0	1	1
4	0	0	1	0

	S_1	S_2	S_3	S_4
h_1	1	3	0	1
h_2	0	2	0	0

	$SIM(S_1, S_2)$	$SIM(S_1, S_3)$	$SIM(S_1, S_4)$
True	0	1/4	2/3
Estimated	0	1/2	1

- The estimates become close as the number of hash functions increases

Locality-Sensitive Hashing

- Minhashing

- Compresses large documents into small signatures, while preserving the expected similarity of any pair of documents

- Still, the number of pairs of documents may be too large

- (ex) 1,000,000 documents $\rightarrow {}_{1,000,000}C_2 \approx 500,000,000,000$ pairs

- Locality-sensitive hashing (LSH)


- Allows us to focus our attention **only** on pairs that are likely to be similar, without investigating every pair

General Approach to LSH

- Big idea
 - Hash items *several times*
 - We call a pair that hashed to the same bucket for any of the hashings to be a *candidate pair*
 - We check *only* the candidate pair for similarity
- The hope is that most of the dissimilar pairs will never hash to the same bucket
 - *False positives* will be only a small fraction of all pairs
 - Those dissimilar pairs that do hash to the same bucket
 - *False negatives* will be only a small fraction of the truly similar pairs
 - Those similar pairs that do not hash to the same bucket under at least one of the hash functions


LSH for Minhash Signatures

- Divide the signature matrix into b bands of r rows



- For each band, hash its portion of each column to k buckets
 - Make k as large as possible (so there are almost no collisions)
- Candidate column pairs** are those that hash to the same bucket for one or more bands

(Ex) LSH for Minhash Signatures



Simplifying Assumption

- There are *enough* buckets that columns are unlikely to hash to the same bucket unless they are *identical* in a particular band
- Hereafter, we assume that “*same* bucket” means “*identical* in that band”
- Observation
 - The more similar two columns are, the more likely they will be identical in some band
 - Thus, the banding strategy makes similar columns much more likely to be candidate pairs than dissimilar pairs

Analysis of the Banding Technique (1/2)


- Let (d_1, d_2) be a pair of documents that have Jaccard similarity s
- Can we compute the **probability** that (d_1, d_2) becomes a candidate pair when we use LSH?
 - Obviously, this probability depends on their Jaccard similarity s
- Example
 - If the Jaccard similarity between d_1 and d_2 is 0.8, the probability of (d_1, d_2) becoming a candidate pair is 99.965%
 - If the Jaccard similarity between d_1 and d_2 is 0.3, the probability of (d_1, d_2) becoming a candidate pair is 4.74%

Analysis of the Banding Technique (2/2)

- Suppose we use b bands of r rows each
- Let (d_1, d_2) be a pair of documents that have Jaccard similarity s
 - Thus, the probability the minhash signatures for d_1 and d_2 agree in any one particular row of the signature matrix is s
- The probability that (d_1, d_2) becomes a ***candidate pair***
 - $P(\text{the signatures agree in all rows of one particular band}) = s^r$
 - $P(\text{the signatures disagree in at least one row of a particular band}) = 1 - s^r$
 - $P(\text{the signatures disagree in at least one row of each of the bands}) = (1 - s^r)^b$
 - $P(\text{the signatures agree in all the rows of at least one band}) = 1 - (1 - s^r)^b$
 - This is the probability that (d_1, d_2) becomes a candidate pair

$$1 - (1 - s^r)^b$$

- This function has the form of an *S-curve*
 - Exactly the shape we want (i.e., pairs with similarity above threshold have a high probability of becoming a candidate, while those below the threshold have a low probability of becoming a candidate)



Threshold

- The value of similarity at which the probability of becoming a candidate is $1/2$
 - Roughly where the rise is the steepest
 - Approximately $(1/b)^{1/r}$
- For large b and r
 - Pairs with similarity above the threshold are ***very likely*** to become candidates
 - Pairs with similarity below the threshold are ***unlikely*** to become candidates

(Ex) $1 - (1 - s^r)^b$

- When $b = 20$ and $r = 5$
 - Thus, the length of a signature = 100

- $1 - (1 - s^5)^{20}$


s	$1 - (1 - s^r)^b$
.2	.006
.3	.047
.4	.186
.5	.470
.6	.802
.7	.975
.8	.9996

The probability that their signatures are identical in a particular band = $0.8^5 = 33\%$

- The threshold is just slightly more than 0.5
- If $s = 0.8$, the probability that their signatures are hashed into the same bucket for at least one of 20 bands is 0.9996


Picking r and b (1/4)

- Ideal curve



Picking r and b (2/4)

- When $b = 1$ and $r = 1$




Picking r and b (3/4)

- When b and r are large



Picking r and b (4/4)

- False negative and false positive
 - When $r = 5$ and $b = 10$



Blue area: False Negative rate
Green area: False Positive rate

Combining the Techniques (1/2)

✓ Note that false negatives or false positives can be produced

Step 1: Pick a value of k and construct from each document the set of k -shingles

Step 2: Pick a length n for the minhash signatures and compute the minhash signatures for all the documents

Step 3: Choose a threshold t that defines how similar documents have to be regarded as a similar pair. Pick a number of bands b and a number of rows r such that $br = n$, and the threshold t is approximately $(1/b)^{1/r}$

Combining the Techniques (2/2)

Step 5: Construct candidate pairs by applying the LSH technique

Step 6: Examine each candidate pair's signatures and determine whether the fraction of components in which they agree is at least t

Step 7 (optional): If the signatures are sufficiently similar, go to the document themselves and check that they are truly similar

✓ Documents can have similar signatures by luck

Distance Measures

Distance Measures

- Measure of *closeness*
- Example: Jaccard distance
 - 1 minus the Jaccard similarity
 - A distance measure for *sets*
 - The closer sets are, the lower the Jaccard distance
- There are a number of other distance measures that make sense in some applications
 - Euclidean distance, cosine distance, edit distance, hamming distance

Definition of a Distance Measure

- Suppose we have a set of points, called a *space*
- Let x and y be two points in the space
- A *distance* measure on this space is a function $d(x, y)$ that satisfies the following condition:
 - $d(x, y) \geq 0$ (no negative distance)
 - $d(x, y) = 0$ if and only if $x = y$ (distance is zero from a point to itself)
 - $d(x, y) = d(y, x)$ (distance is symmetric)
 - $d(x, y) \leq d(x, z) + d(z, y)$ (the *triangle inequality*)
 - To travel from x to y , we cannot obtain any benefit if we are forced to travel via some particular third point z

Euclidean Distance (1/2)

- The most familiar distance measure
 - The one we normally think of as “distance”
- n -dimensional Euclidean space
 - A space where points are vectors of n real numbers (e.g., $[x_1, x_2, \dots, x_n]$)
- L_2 -norm

$$d([x_1, x_2, \dots, x_n], [y_1, y_2, \dots, y_n]) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

- Note that all the requirements for a distance measure are satisfied

Euclidean Distance (2/2)

- L_r -norm (general form)

$$d([x_1, x_2, \dots, x_n], [y_1, y_2, \dots, y_n]) = \left(\sum_{i=1}^n |x_i - y_i|^r \right)^{1/r}$$

- L_1 -norm (called ***Manhattan distance***)

- The distance one would have to travel between points if one were constrained to travel among the streets of a city such as Manhattan

- L_∞ -norm

$$d([x_1, x_2, \dots, x_n], [y_1, y_2, \dots, y_n]) = \max_i |x_i - y_i|$$

- The maximum of $|x_i - y_i|$ over all dimensions i
- As r gets larger, only the dimension with the largest difference matters

(Ex) Euclidean Distance

- Consider the two-dimensional Euclidean space and the two points $(2, 7)$ and $(6, 4)$
- L_2 -norm
 - $(|2 - 6|^2 + |7 - 4|^2)^{1/2} = 5$
- L_1 -norm
 - $|2 - 6| + |7 - 4| = 7$
- L_∞ -norm
 - $\max(|2 - 6|, |7 - 4|) = 4$

Jaccard Distance


- $d(x, y) = 1 - SIM(x, y) = 1 - |x \cap y|/|x \cup y|$
 - In other words, the probability that a random minhash function does **not** send x and y to the same value
- Note that all the requirements for a distance measure are satisfied
 - $d(x, y) \geq 0$, because $|x \cap y| \leq |x \cup y|$
 - $d(x, y) = 0$ iff $x = y$, because $x \cap x = x \cup x = x$
 - $d(x, y) = d(y, x)$, because $x \cup y = y \cup x$ and $x \cap y = y \cap x$
 - $d(x, y) \leq d(x, z) + d(z, y)$
 - We show $P(h(x) \neq h(y)) \leq P(h(x) \neq h(z)) + P(h(z) \neq h(y))$
 - This is true because whenever $h(x) \neq h(y)$, at least one of $h(x)$ and $h(y)$ must be different from $h(z)$

Cosine Distance

- The cosine distance between two points $x = [x_1, x_2, \dots, x_n]$ and $y = [y_1, y_2, \dots, y_n]$ measures the **angle** that the vectors to those points make

$$\cos(\theta) = \frac{x \cdot y}{||x|| \cdot ||y||} = \frac{\sum_{i=1}^n x_i y_i}{\sqrt{\sum_{i=1}^n x_i^2} \sqrt{\sum_{i=1}^n y_i^2}}$$

- This angle will be in the range 0° to 180°



- Note that all the requirements for a distance measure are satisfied


Edit Distance (1/2)

- A distance measure for *strings*
- The edit distance between two strings $x = x_1x_2\dots x_n$ and $y = y_1y_2\dots y_n$ is the ***smallest*** number of insertions and deletions of single characters that will convert x to y
- Example
 - The edit distance between $x = abcde$ and $y = acfdeg$ is 3
 - To convert x to y
 1. Delete b
 2. Insert f after c
 3. Insert g after e
 - No sequence of fewer than 3 insertions and/or deletions will convert x to y

Edit Distance (2/2)

- Calculation of the edit distance $d(x, y)$

- Compute a longest common subsequence (LCS) of x and y
- $d(x, y) = |x| + |y| - 2|\text{LCS of } x \text{ and } y|$



- Example

- $d(\text{abcde}, \text{acfddeg}) = |\text{abcde}| + |\text{acfddeg}| - 2|\text{acde}| = 5 + 6 - 2 \cdot 4 = 3$
- $d(\text{aba}, \text{bab}) = |\text{aba}| + |\text{bab}| - 2|\text{ab}| = 3 + 3 - 2 \cdot 2 = 2$
 - Or, $|\text{aba}| + |\text{bab}| - 2|\text{ba}| = 3 + 3 - 2 \cdot 2 = 2$

- Note that all the requirements for a distance measure are satisfied

Hamming Distance

- A distance measure for *vectors*
- The Hamming distance between two vectors is the number of components in which they *differ*
 - (ex) $d([2, 1, 7], [2, 2, 3]) = 2$, $d(10101, 11110) = 3$
- Note that all the requirements for a distance measure are satisfied
 - $d(x, y) \geq 0$ (it is clear)
 - $d(x, y) = 0$ iff $x = y$ (if the distance is zero, then the vectors are identical)
 - $d(x, y) = d(y, x)$ (the distance doesn't depend on the order of two vectors)
 - $d(x, y) \leq d(x, z) + d(z, y)$
 - If x and z differ in m components, and z and y differ in n components, then x and y cannot differ in more than $m + n$ components