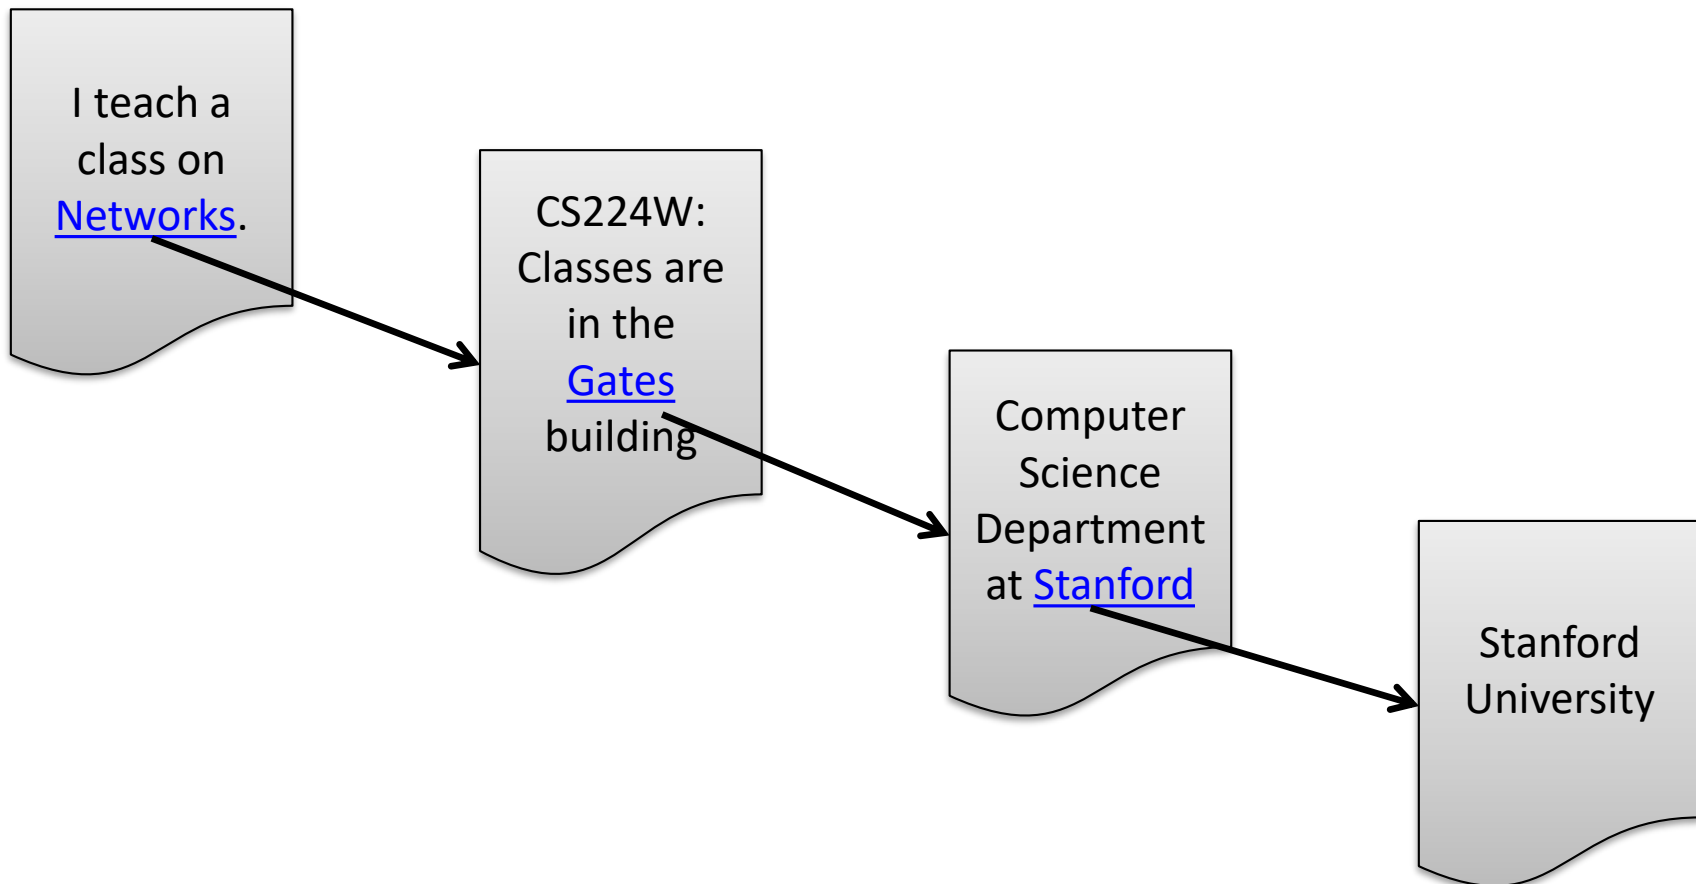

Chapter 5

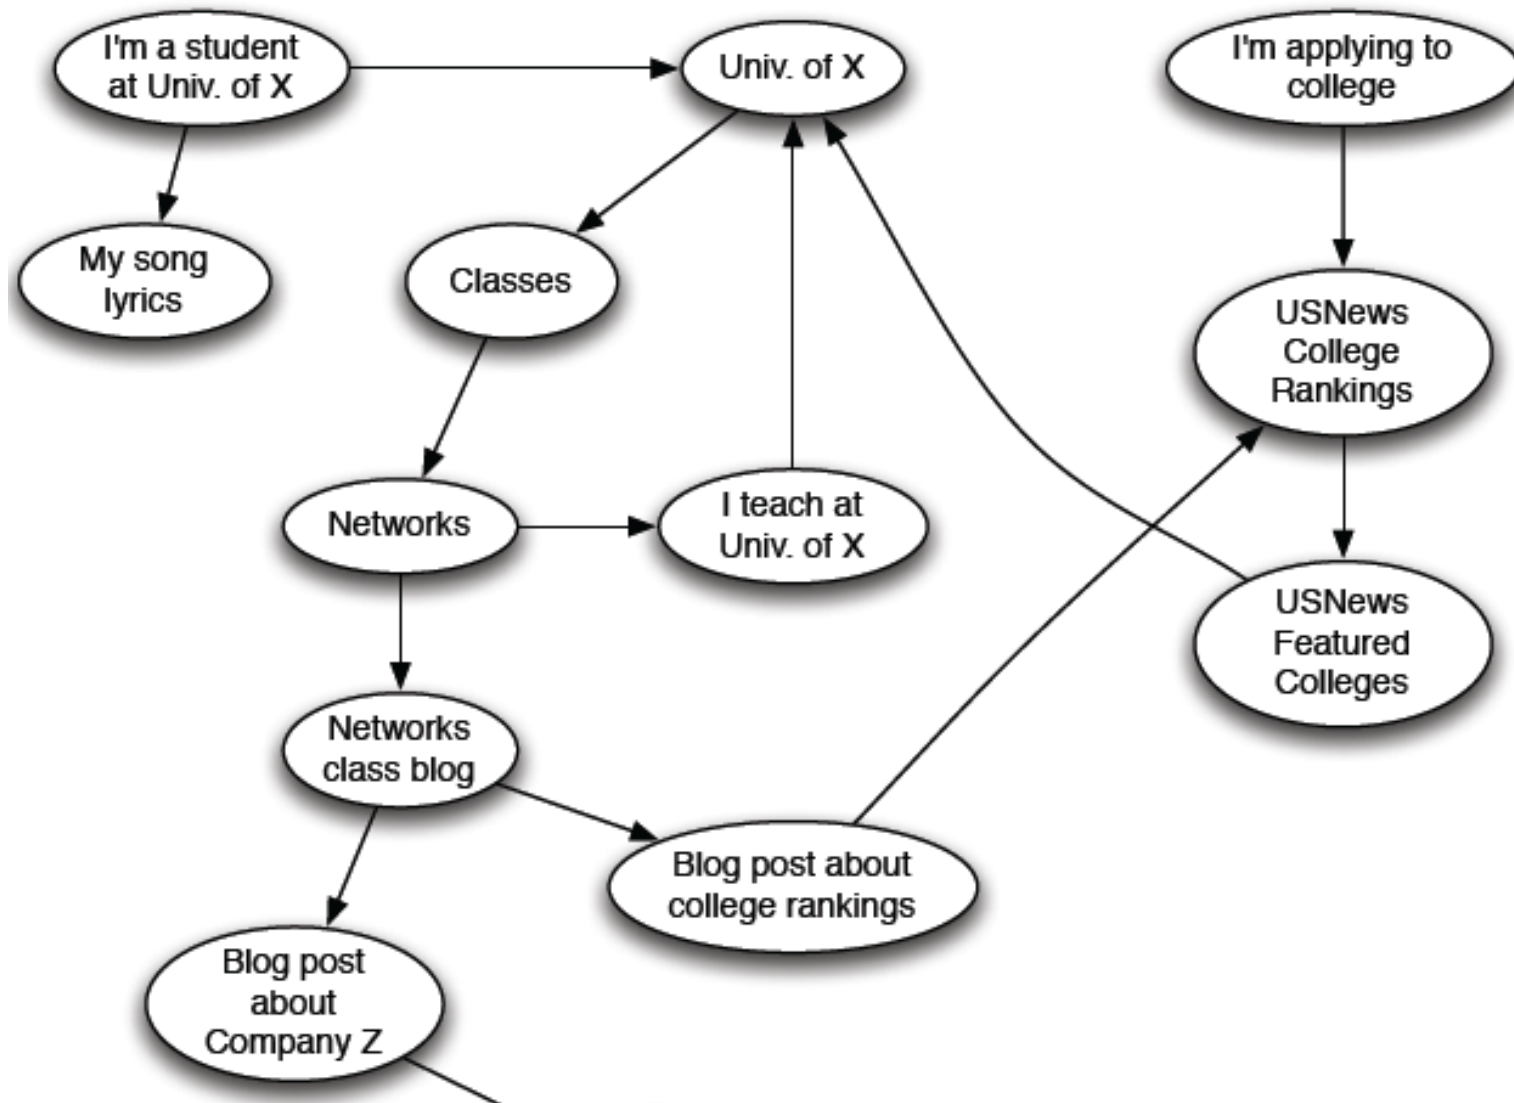
Link Analysis

Web as a Graph

- We can see the Web as a *directed graph*
 - Nodes: Web pages
 - Edges: Hyperlinks



Web as a Directed Graph



How Can We Organize the Web?

- **First try:** Human curated Web directories

- Yahoo, DMOZ, LookSmart

- **Second try:** Web search

- Find relevant Web pages automatically
- (c.f.) Information retrieval

- Find relevant documents in a small and trusted set
- (ex) Newspaper articles, patents, etc.

- However, the Web is *huge*, full of *untrusted* documents, random things, Web spam, etc.



Web Search: Two Challenges

(1) Web contains many sources of information

- Who to **trust**?
- Trick: Trustworthy pages may point to each other!


(2) What is the **best** answer to query “newspaper”?

- No single right answer
- Trick: Pages that actually know about newspapers might all be pointing to many newspapers



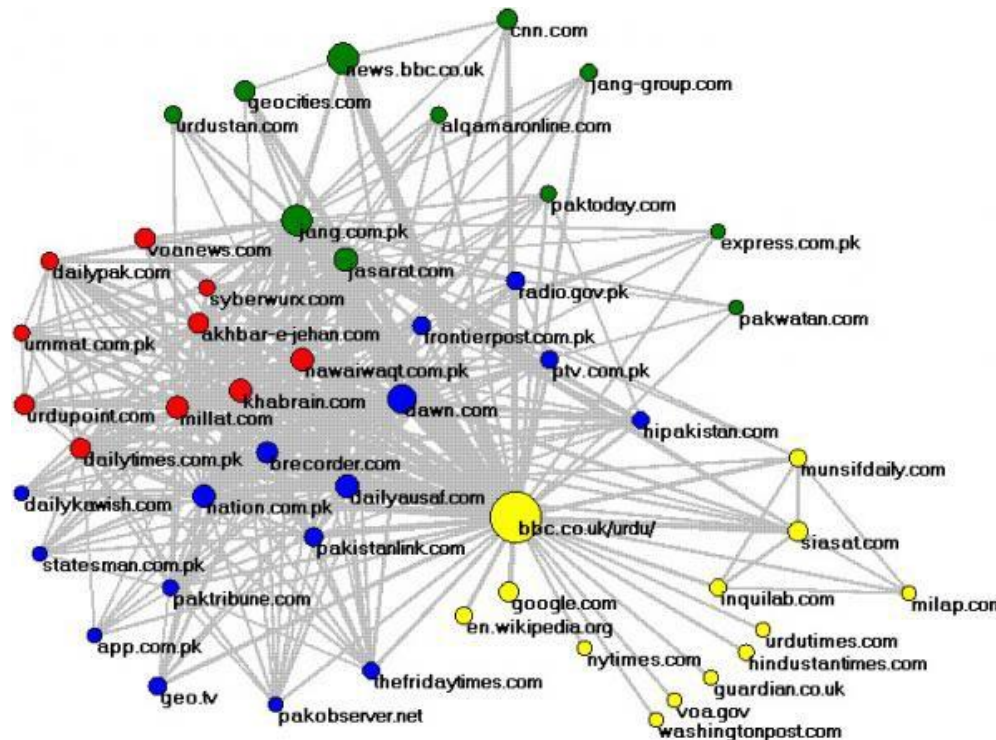
Ranking Nodes on the Graph

- All web pages are not equally “*important*”
 - (ex) www.joe-schmoe.com (1 in-link) vs. www.stanford.edu (23,400 in-links)
- There is large diversity in the web-graph node connectivity
- Let’s *rank* the pages by the link structure!



Link Analysis Algorithms

- We will cover the following *Link Analysis approaches* for computing *importances* of nodes in a graph:
 - Page Rank
 - Topic-Specific (Personalized) Page Rank
 - Web Spam Detection Algorithms



PageRank

Web Search

- *Efficient* and *accurate* Web search has changed our lives greatly
 - Through search engines such as Google
- Google “PageRank”
 - The first able to defeat *spammers* who had made search almost useless
 - We shall explain what it is and how it is computed efficiently
- Variations on PageRank
 - **TrustRank**: prevents spammers attacking PageRank (called link spam)
 - **Topic-sensitive PageRank**: weights Web pages based on their topic
 - **HITS**: “hubs and authorities” approach to evaluating Web pages

Early Search Engines Before Google

- Crawl the Web and list the terms found in each page
- Construct an inverted index
 - A data structure used for finding all the places where a given term occurs
- When a search query (list of terms) is issued
 - Use the inverted index to retrieve the pages with those terms
 - **Rank** them in a way that reflects the use of the terms within the page
- Factors for page ranking
 - The number of occurrences of the term in the page
 - The place of occurrence of the term in the page (header? body?)

Term Spam

- As people began to use search engines, unethical people tried to *fool* search engines into leading people to their page
- Examples
 - Add a term to your page thousands of times
 - A search engine would think your page is very important about the term
 - Thus, the search engine would list your page first
 - Make them the same color as background to hide their occurrences
- *Term spam*
 - Techniques for fooling search engines into believing your page is about something is not

Google's Innovations to Combat Term Spam

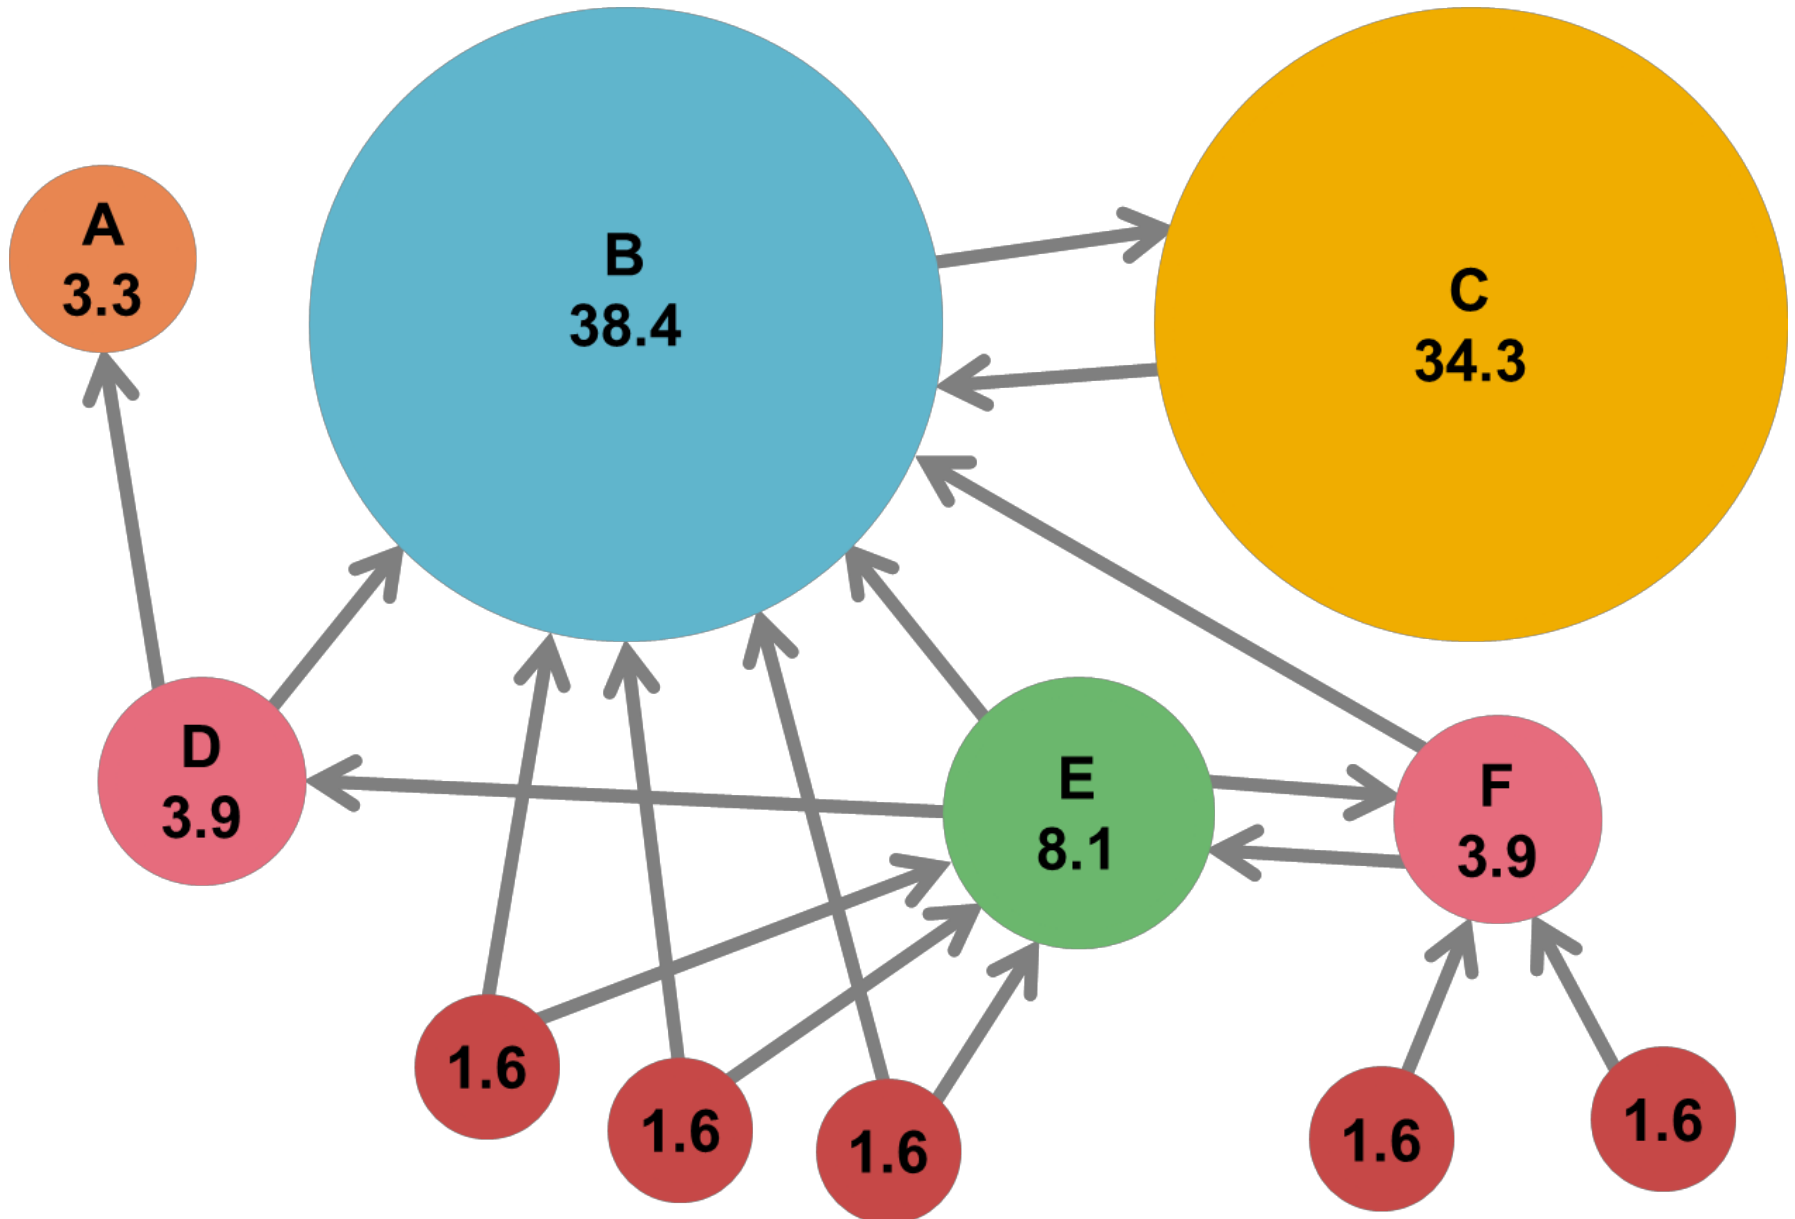
1. PageRank was used to simulate Web surfers

- They start at a random page
- They follow randomly chosen outlinks from their current page
- They would tend to congregate if this process iterates many time
- Pages with a large number of surfers are considered more “**important**”
- Google prefers important pages when deciding which pages show first

2. The content of a page was judged not only by the terms appearing on that page, but by the terms used in or near ***the links to that page***

- Note that it is **not** easy for a spammer to add false terms to a page they do not control

Example: PageRank Scores



Why The Two Techniques Work? (1/2)

- Google believes what other pages say about him, over what he says about himself
 - While a spammer can still add a false term to his page, the above fact would negate the use of false terms
- What if the spammer creates many pages of his own, and links to his page with a link with the false term?
 - But those pages would not be given much importance by PageRank, since other pages would not link to them
 - Therefore, he still would not be able to fool Google into thinking this page is important

Why The Two Techniques Work? (2/2)

- Why should simulation of random surfers allow us to approximate the intuitive notion of the “*importance*” of pages?
- Two related motivations
 - Users of the Web “vote with their feet”
 - They tend to place links to pages they think are good or useful pages to look at, rather than bad or useless pages
 - The behavior of a random surfer indicates which pages users of the Web are likely to visit
 - Users are more likely to visit useful pages than useless pages
- But regardless of the reason, the PageRank measure has been proved *empirically* to work!


Definition of PageRank

- PageRank

- A function that assigns a real number to each page in the Web
- The higher the PageRank of a page, the more “*important*” it is



- PageRank represents the Web as a *direct graph*

- Pages are the nodes
- There is an arc from page p_1 to page p_2 if there are one or more links from p_1 to p_2



Transition Matrix of the Web

- Describes what happens to random surfers after one step
- The transition matrix M is an $n \times n$ matrix
 - n : the number of pages
 - The element m_{ij} in row i and column j
 - $1/k$, if page j has k arcs out, and one of them is to page i
 - 0, otherwise
 - The **probability** that a surfer at page i will move to page j at the next step



$$M = \begin{bmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix} \begin{matrix} A \\ B \\ C \\ D \end{matrix}$$

The Probability Distribution Vector

- The probability distribution vector \mathbf{v}
 - whose j th component is the probability that the surfer is at page j
 - v_j is the (idealized) **PageRank** function of page j
 - Initially, each component has an equal value of $1/n$
- The distribution vector \mathbf{v} after one step
 - $\mathbf{v} \leftarrow M\mathbf{v}$
 - Because the probability that a random surfer will be at page i at the next step is $\sum_j m_{ij}v_j$
- The distribution vector \mathbf{v} after two steps
 - $\mathbf{v} \leftarrow M(M\mathbf{v}) = M^2\mathbf{v}$

A Limiting Distribution of \mathbf{v}

- This sort of behavior is an example of Markov process
- It is known that the distribution of the surfer approaches a *limiting distribution* \mathbf{v} that satisfies

$$\mathbf{v} = M\mathbf{v}$$

provided two conditions are met:

- The graph is *strongly connected*
 - That is, it is possible to get from any node to any other node
- There are no *dead ends*
 - Nodes that have no arc out

Computing a Limiting Distribution of \mathbf{v}

- The limit is reached when multiplying \mathbf{v} by M another time does not change \mathbf{v}
 - i.e., $\mathbf{v} = M\mathbf{v}$
- We can compute the limiting distribution \mathbf{v} by
 - Starting with the initial vector \mathbf{v}_0 and
 - Multiplying by M some number of times, until the vector we get shows little change at each round
 - (ex) $\mathbf{v}_1 = M\mathbf{v}_0 \rightarrow \mathbf{v}_2 = M\mathbf{v}_1 \rightarrow \mathbf{v}_3 = M\mathbf{v}_2 \rightarrow \dots$ until $\mathbf{v}_{i+1} \cong \mathbf{v}_i$
- In practice, for the Web itself, 50-75 iterations are sufficient to **converge** to within the error limits of double-precision arithmetic

Example


- Consider the matrix M in the previous example
 - Because there are 4 nodes, the initial vector $\mathbf{v}_0 = [1/4, 1/4, 1/4, 1/4]^T$
- The sequence of approximations to the limit that we get by multiplying at each step by M

$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix}, \begin{bmatrix} 9/24 \\ 5/24 \\ 5/24 \\ 5/24 \end{bmatrix}, \begin{bmatrix} 15/48 \\ 11/48 \\ 11/48 \\ 11/48 \end{bmatrix}, \begin{bmatrix} 11/32 \\ 7/32 \\ 7/32 \\ 7/32 \end{bmatrix}, \dots, \begin{bmatrix} 3/9 \\ 2/9 \\ 2/9 \\ 2/9 \end{bmatrix}$$

- The probability that a surfer being at A is larger than that for B, C, D
 - Therefore, A is a more **important** page than B, C, D
- In the real Web, the true probability of being at a node like www.amazon.com is orders of magnitude greater than the probability of typical nodes

Structure of the Web

- The “bowtie” picture of the Web
 - In practice, the Web is **not** strongly connected



Large Components

- A large strongly connected component (SCC)
- In-component
 - Pages that could reach the SCC by following links, but not reachable from the SCC
- Out-component
 - Pages reachable from the SCC but unable to reach the SCC
- Tendrils
 - (Type 1) Pages reachable from the in-component but not able to reach the in-component
 - (Type 2) Pages that can reach the out-component but not reachable from the out-component

Small Components

- Tubes

- Pages reachable from the in-component and able to reach the out-component, but unable to the SCC or be reached from the SCC

- Isolated component

- Pages unreachable from the large components (the SCC, in- and out-components) and unable to reach those components

Problems with the Structure of the Web


- Several of these structures *violate* the assumption needed for the Markov-process iteration to converge to a limit
- Example
 - When a random surfer enters either the out-component or a tendrill off the in-component, they can *never* leave
 - Thus, *no* page in the SCC or in-component winds up with any probability of a surfer being there
 - Consequently, we conclude *falsely* that nothing in the SCC or in-component is of any importance

How PageRank Prevent Such Anomalies

- There are really *two problems* we need to avoid

① Dead end

- A page that has no links out
- No page that can reach a dead end can have any PageRank at all



② Spider trap


- A group of pages that all have outlinks but never link to any other pages
- ✓ Both these problems are solved by a method called “*taxation*”
- A random surfer has a finite probability of leaving the Web at any step
 - New surfers are started at each page

Avoiding Dead Ends (1/2)

- If we allow dead ends, then some of the columns of the transition matrix will sum to 0 rather than 1
 - Consequently, some or all of the components of $M^i \mathbf{v}$ go to 0
 - That is, importance “drains out” of the Web
 - We get no information about the relative importance of pages

- Example

- Consider the following transition matrix, where C is a dead end
 - The sum of the third column, for C , is 0, not 1



$$M = \begin{bmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix} \begin{matrix} A \\ B \\ C \\ D \end{matrix}$$

Avoiding Dead Ends (2/2)

- Example (cont'd)

- The sequence of vectors that result by starting the vector with each component $1/4$, and repeatedly multiplying the vector by M

$$M = \begin{bmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix}, \begin{bmatrix} 3/24 \\ 5/24 \\ 5/24 \\ 5/24 \end{bmatrix}, \begin{bmatrix} 5/48 \\ 7/48 \\ 7/48 \\ 7/48 \end{bmatrix}, \begin{bmatrix} 21/288 \\ 31/288 \\ 31/288 \\ 31/288 \end{bmatrix}, \dots, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

- As we see, the probability of a surfer being anywhere goes to 0, as the number of steps increase

Two Approaches to Dealing with Dead Ends

1. Recursive deletion of dead ends

- Drop the dead ends from the graph, and also drop their incoming arcs
- Doing so may create more dead ends, which also have to be dropped, recursively
- However, eventually we wind up with a strongly-connected component, none of whose nodes are dead ends

2. Taxation

- Modify the process by which random surfers are assumed to move about the Web
- This method also solves the problem of spider traps

Recursive Deletion of Dead Ends (1/2)


1. Drop the dead ends from the graph recursively
2. Solve the remaining graph G
 - By whatever means are appropriate, including the taxation method
3. Restore the graph
 - But keep the PageRank values for the nodes of G
4. Compute the PageRank of nodes not in G , but with predecessors all in G
 - By summing, over all predecessors p , the PageRank of p divided by the number of successors of p in the full graph

Recursive Deletion of Dead Ends (2/2)

5. Compute the PageRank of other nodes, not in G , that have the PageRank of all their predecessors computed
 - By the same process
6. Eventually, all nodes outside G will have their PageRank computed
 - They can surely be computed in the order ***opposite*** to that in which they were deleted



Example (1/3)

- Consider the following graph, where E is a dead end
 - After removing E , we find that C is now a dead end
 - After removing C , there are no more dead ends



Example (2/3)

- The transition matrix M for the remaining graph G



$$M = \begin{matrix} & \begin{matrix} A & B & D \end{matrix} \\ \begin{bmatrix} 0 & 1/2 & 0 \\ 1/2 & 0 & 1 \\ 1/2 & 1/2 & 0 \end{bmatrix} & \begin{matrix} A \\ B \\ D \end{matrix} \end{matrix}$$

- We get the PageRanks for this matrix


- By starting with a vector with all components equals to $1/3$, and repeatedly multiplying by M

$$\begin{bmatrix} 1/3 \\ 1/3 \\ 1/3 \end{bmatrix}, \begin{bmatrix} 1/6 \\ 3/6 \\ 2/6 \end{bmatrix}, \begin{bmatrix} 3/12 \\ 5/12 \\ 4/12 \end{bmatrix}, \begin{bmatrix} 5/24 \\ 11/24 \\ 8/24 \end{bmatrix}, \dots, \boxed{\begin{bmatrix} 2/9 \\ 4/9 \\ 3/9 \end{bmatrix}} \begin{matrix} A \\ B \\ D \end{matrix}$$

Example (3/3)


- Now we compute PageRank for C

- Because all its predecessors (i.e., A and B) have PageRanks computed
- The PageRank for $C = 1/3 \times 2/9 + 1/2 \times 3/9 = 13/54$



- Finally we compute PageRank for E

- Because all its predecessors (i.e., C) have PageRanks computed
- The PageRank for $E = 13/54$




- ✓ Note that the sums of the PageRank **exceed** 1

- They no longer represent the distribution of a random surfer
- Yet they represent decent estimates of the relative importance of the pages



Spider Trap

- A set of nodes with no dead ends but no arcs out
- Can appear *intentionally* or unintentionally on the Web
- Causes the PageRank calculation to place all the PageRank within the spider trap



Example

- Consider the following graph, where C is a spider trap of 1 node, and its transition matrix M



$$M = \begin{bmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 1 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$

- The usual iteration to compute the PageRank

$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix}, \begin{bmatrix} 3/24 \\ 5/24 \\ 11/24 \\ 5/24 \end{bmatrix}, \begin{bmatrix} 5/48 \\ 7/48 \\ 29/48 \\ 7/48 \end{bmatrix}, \begin{bmatrix} 21/288 \\ 31/288 \\ 205/288 \\ 31/288 \end{bmatrix}, \dots, \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

- As predicted, ***all the PageRank is at C***
- Since once there a random surfer can never leave

Taxation

- To avoid this problem, we allow each random surfer a small probability of *teleporting* to a random page
 - Rather than following an out-link from their current page
- The calculation of PageRank is modified as follows:

$$\mathbf{v}' = \beta M \mathbf{v} + (1 - \beta) \mathbf{e}/n$$



- β : a chosen constant (usually in the range 0.8 to 0.9)
- \mathbf{e} : a vector of all 1's with the appropriate number of components
- n : the number of nodes in the Web graph
- $\beta M \mathbf{v}$: represents the case where, with probability β , the random surfer follows an out-link from their present page
- $(1 - \beta) \mathbf{e}/n$: the introduction, with probability $1 - \beta$, of a new random surfer at a random page

Taxation and Dead Ends

- If the graph has no dead ends
 - The surfer decides either to follow a link or teleport to a random page
- If the graph has dead ends
 - There is a third possibility that the surfer goes nowhere
 - Due to the term $(1 - \beta)c/n$, there will always be some fraction of a surfer operating on the Web
 - Thus, when there are dead ends, the sum of the components of \mathbf{v} may less than 1, but it will **never** reach 0

Example (1/2)

- Suppose we apply the new approach to the following graph



$$M = \begin{bmatrix} 0 & 1/2 & 0 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 1 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$

- If we use $\beta = 0.8$, then the equation for the iteration becomes

$$\mathbf{v}' = \begin{bmatrix} 0 & 2/5 & 0 & 0 \\ 4/15 & 0 & 0 & 2/5 \\ 4/15 & 0 & 4/5 & 2/5 \\ 4/15 & 2/5 & 0 & 0 \end{bmatrix} \mathbf{v} + \begin{bmatrix} 1/20 \\ 1/20 \\ 1/20 \\ 1/20 \end{bmatrix}$$

Example (2/2)

- Here are the first few iterations:

$$\begin{bmatrix} 1/4 \\ 1/4 \\ 1/4 \\ 1/4 \end{bmatrix}, \begin{bmatrix} 9/60 \\ 13/60 \\ 25/60 \\ 13/60 \end{bmatrix}, \begin{bmatrix} 41/300 \\ 53/300 \\ 153/300 \\ 53/300 \end{bmatrix}, \begin{bmatrix} 543/4500 \\ 707/4500 \\ 2543/4500 \\ 707/4500 \end{bmatrix}, \dots, \begin{bmatrix} 15/148 \\ 19/148 \\ \boxed{95/148} \\ 19/148 \end{bmatrix}$$

- By being a spider trap, C has managed to get more than half of the PageRank for itself
- However, the effect has been *limited*, and each of the nodes gets some of the PageRank

Using PageRank in a Search Engine (1/2)

- After a search engine has crawled the portion of the Web, the ***PageRank vector*** for that portion is calculated
- Each search engine has a ***secrete formula*** that decides the order in which to show pages to the user in response to a search query consisting of one or more search terms (words)
- Google is said to use over 250 different properties of pages, from which a linear order of pages is decided

Using PageRank in a Search Engine (2/2)

- First, in order to be considered for the ranking at all, a page has to have at least one of the search terms in the query
 - Normally, a page has very little chance of being in the top ten unless all the search terms are present
- Among the qualified pages, a score is computed for each
 - An important component of this score is the **PageRank** of the page
 - Other components include the presence or absence of search terms in prominent places, such as headers or the links to the page itself

Efficient Computation of PageRank

Efficient Computation of PageRank

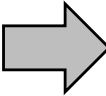
- To compute the PageRank, we have to perform a matrix-vector multiplication on the order of **50** times
 - Until the vector is close to unchanged at one iteration
- To a first approximation, the MapReduce method given in Chapter 2 is suitable
- However, we must deal with two issues:
 1. The transition matrix of the Web M is very ***sparse***
 - Representing it by all its elements is highly inefficient
 2. We may ***not*** be using MapReduce, or may wish to use a ***combiner***
 - In this case, the striping approach discussed in Chapter 2 is not sufficient to avoid heavy use of disk (thrashing)

Representing Transition Matrices

- The transition matrix is very *sparse*
 - Since the average Web page has about 10 out-links
 - (ex) If we are analyzing a graph of 10 billion pages, then only one in a billion entries is not 0
- The proper way to represent a sparse matrix
 - List the locations of the **nonzero** entries and their values
 - Thus, the space needed is linear in the number of nonzero entries, rather than quadratic in the side of the matrix
 - (ex) 16 bytes per nonzero entry (4-byte integers for coordinates of an elements and an 8 byte double-precision number for the value)

Further Compression for a Transition Matrix

- If we list the nonzero entries by columns, then we know what the value of each nonzero entry is
 - It is 1 divided by the out-degree of the page
- We can thus represent a column by
 - One integer for the out-degree
 - One integer per nonzero entry in that column (i.e., the row number)
- Example

$$M = \begin{bmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$


Source	Degree	Destinations
<i>A</i>	3	<i>B, C, D</i>
<i>B</i>	2	<i>A, D</i>
<i>C</i>	1	<i>A</i>
<i>D</i>	2	<i>B, C</i>

A compact representation of M

PageRank Iteration Using MapReduce

- Recall that each iteration of the PageRank algorithm computes

$$\mathbf{v}' = \beta M \mathbf{v} + (1 - \beta) \mathbf{e}/n$$


- If each Map task can store the full vector \mathbf{v} in memory and also have room in main memory for the result vector \mathbf{v}'
 - Then there is little more here than a matrix-vector multiplication
 - Additional steps: to multiply each component of $M \mathbf{v}$ by constant β and to add $(1 - \beta)/n$ to each component
- However, it is likely that \mathbf{v} is much too large to fit in main memory
 - Then, we can use the striping method described in Chapter 2 to execute the MapReduce process efficiently
 - That is, we break M into vertical stripes and break \mathbf{v} into corresponding horizontal stripes

Use of Combiners

- The previous method might **not** be adequate for two reasons:
 1. We might wish to add terms for \mathbf{v}_i' at the Map tasks
 - \mathbf{v}_i' : the i th component of the result vector \mathbf{v}
 - The Reduce function simply adds terms with a common key
 - Recall that for a MapReduce implementation of matrix-vector multiplication, the key is the value of i for which a term $m_{ij}\mathbf{v}_j$ is intended
 2. We might not be using MapReduce at all
- We are trying to implement a **combiner** in conjunction with a Map task
 - The second case uses essentially the same idea

An Alternative Strategy (1/2)

- Partition the matrix into k^2 blocks
 - The vectors are still partitioned into k stripes
 - (ex) partitioning the matrix with $k = 4$



- We use k^2 Map tasks
 - Each task gets one square of M , say M_{ij} , and one stripe of \mathbf{v} , \mathbf{v}_j
 - Each stripe of \mathbf{v} is sent to k different Map tasks
 - $\mathbf{v}_j \rightarrow M_{1j}, M_{2j}, \dots, M_{kj}$

An Alternative Strategy (2/2)

■ The transmission cost

- Each \mathbf{v}_j is transmitted over the network k times
- Each M_{ij} is transmitted only once
- Since the size of M is expected to be several times the size of \mathbf{v} , the transmission cost is not too much greater than the minimum possible
- Because we are doing considerable combining at the Map tasks, we save as data is passed from the Map tasks to the Reduce tasks


■ The advantage of this approach

- We can keep both the j th stripe of \mathbf{v} and the i th stripe of \mathbf{v}' in main memory as we process M_{ij}
- Thus, we don't need to access disk frequently to bring them into main memory (***thrashing***)

Representing Blocks of the Transition Matrix


- Since we are representing transition matrices in the special way described previously, we need to consider how the blocks of a transition matrix are represented
- Example
 - Suppose the following matrix is partitioned into blocks with $k = 2$

$$M = \begin{bmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix}$$



$$M = \left[\begin{array}{cc|cc} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ \hline 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{array} \right]$$

Source	Degree	Destinations
<i>A</i>	3	<i>B, C, D</i>
<i>B</i>	2	<i>A, D</i>
<i>C</i>	1	<i>A</i>
<i>D</i>	2	<i>B, C</i>



?

Representing Blocks of the Transition Matrix

- Example (cont'd)

$$M = \begin{matrix} & \begin{matrix} A & B & C & D \end{matrix} \\ \begin{bmatrix} 0 & 1/2 & 1 & 0 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 0 & 0 & 1/2 \\ 1/3 & 1/2 & 0 & 0 \end{bmatrix} & \begin{matrix} A \\ B \\ C \\ D \end{matrix} \end{matrix}$$

Source	Degree	Destinations
<i>A</i>	3	<i>B</i>
<i>B</i>	2	<i>A</i>

Representation of M_{11}
(connecting *A* and *B* to *A* and *B*)

Source	Degree	Destinations
<i>A</i>	3	<i>C, D</i>
<i>B</i>	2	<i>D</i>

Representation of M_{21}
(connecting *A* and *B* to *C* and *D*)

Source	Degree	Destinations
<i>C</i>	1	<i>A</i>
<i>D</i>	2	<i>B</i>

Representation of M_{12}
(connecting *C* and *D* to *A* and *B*)

Source	Degree	Destinations
<i>D</i>	2	<i>C</i>

(we can avoid the entry for *C*)


Representation of M_{22}
(connecting *C* and *D* to *C* and *D*)

Other Efficient Approaches to PageRank Iteration

- The k^2 partitioning algorithm is not the only option
 - There are several other approaches that use ***fewer*** processors
- These algorithms share the following good property
 - The matrix M is read only once, although the vector \mathbf{v} is read k times
 - k is chosen so that $1/k$ th of \mathbf{v} and \mathbf{v}' can be held in main memory
- Recall that the k^2 partitioning algorithm uses k^2 processors, assuming all Map tasks are executed in parallel at different processors

Other Efficient Approaches to PageRank Iteration

- We can assign all the blocks in one row to a single Map task
 - Thus the number of Map tasks is reduced to k



- We can read the blocks one-at-a-time, so the matrix does not consume a significant amount of main memory
- At the same time that we read M_{ij} , we must read the vector stripe \mathbf{v}_j
- As a result, each of the k Map tasks reads the entire vector \mathbf{v}
- **Advantage:** each Map task can combine *all* the terms for the portion \mathbf{v}_i' for which it is exclusively responsible

Other Efficient Approaches to PageRank Iteration

- We can extend this idea to an environment where MapReduce is *not* used
 - Suppose we have a single processor, with M and \mathbf{v} stored on its disk
- We simulate each Map task one by one
 - The first task: reads M_{11} through M_{1k} and all of \mathbf{v} to compute \mathbf{v}_1'
 - The second task: reads M_{21} through M_{2k} and all of \mathbf{v} to compute \mathbf{v}_2'
 - ...
 - The k th task: reads M_{k1} through M_{kk} and all of \mathbf{v} to compute \mathbf{v}_k'
- We can make k as small as possible
 - Subject to the constraint that there is enough main memory to store $1/k$ th of \mathbf{v} and $1/k$ th of \mathbf{v}' , along with as small a portion of M as we can read from disk

Some Problems with PageRank

- Measures *generic* popularity of a page
 - Biased against topic-specific authorities
 - **Solution:** Topic-Specific PageRank (next)
- Susceptible to *Link spam*
 - Artificial link topographies created in order to boost page rank
 - **Solution:** TrustRank
- Uses a *single* measure of importance
 - Other models of importance
 - **Solution:** Hubs-and-Authorities