

---

# Chapter 6

## Frequent Itemsets


---

# Market-Basket Model

# The Discovery of Frequent Itemsets

---

- Often viewed as the discovery of “*association rules*”
  - Although the latter is a more complex characterization of data
  - Because it depends fundamentally on the discovery of frequent itemsets
- The “*market-basket*” model of data
  - Essentially a many-to-many relationship between “items” and “baskets”



- The “*frequent-itemsets*” problem
  - Find sets of items that appear in (are related to) many of the same baskets

# Algorithms for Finding Frequent Itemsets

---


- A-Priori algorithm
  - Eliminates most large sets as candidates by looking first at smaller sets and recognizing that ***a large set cannot be frequent unless all its subsets are***
- Various improvements to the basic A-Priori idea
  - Concentrating on ***very large data sets*** that stress the available main memory
- Approximate algorithms
  - Work faster but are not guaranteed to find all frequent itemsets
  - Also in this class are those that exploit parallelism (e.g., MapReduce)

# The Market-Basket Model

- Used to describe a common form of many-to-many relationship between two kinds of objects
  - **Items** (e.g., things sold in a supermarket)
  - **Baskets** (sometimes called “transactions”)
  - Each basket consists of a set of items (an **itemset**)
    - Usually we assume that the number of times in a basket is **small**
  - The number of baskets is very **large**
    - Bigger than what can fit in main memory

<i>TID</i>	<i>Items</i>
1	Bread, Coke, Milk
2	Beer, Bread
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Coke, Diaper, Milk

Input



**Association rules discovered:**

$\{\text{Milk}\} \rightarrow \{\text{Coke}\}$

$\{\text{Diaper, Milk}\} \rightarrow \{\text{Beer}\}$

Output

# Definition of Frequent Itemsets


- Frequent itemsets

- A set of items that appears in *many* baskets

- Formal definition

- The **support** for an itemset  $I$ : the number of baskets for which  $I$  is a subset
- We are given a support threshold  $s$
- We say  $I$  is **frequent** if its support  $\geq s$

<i>TID</i>	<i>Items</i>
1	Bread, Coke, Milk
2	Beer, Bread
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Coke, Diaper, Milk



The support for {Beer, Bread} = 2

# (Ex) Frequent Itemsets

---

- Items = {milk, coke, pepsi, beer, juice}
- Support threshold = 3

$$B_1 = \{m, c, b\}$$

$$B_2 = \{m, p, j\}$$

$$B_3 = \{m, b\}$$

$$B_4 = \{c, j\}$$

$$B_5 = \{m, p, b\}$$

$$B_6 = \{m, c, b, j\}$$

$$B_7 = \{c, b, j\}$$

$$B_8 = \{b, c\}$$

- Frequent itemsets (i.e., support  $\geq 3$ )

$\{m\}, \{c\}, \{b\}, \{j\}, \{m, b\}, \{b, c\}, \{c, j\}$

# Applications of Frequent Itemsets (1/3)

---

- The original application: analysis of true market baskets
  - **Items:** the different products that the store sells
  - **Baskets:** the sets of items in a single market basket
  - **Application**
    - A major chain might sell  $10^5$  different items and have millions of baskets
    - By finding frequent itemsets, a retailer can learn what is commonly bought together
    - Especially important are the *unexpected* sets of items bought together
      - (ex) {Bread, Milk} (X), {Diaper, Beer} (O)
    - It offers the supermarket an opportunity to do some clever marketing
      - (ex) Advertise a sale on diapers and raise the price of beer
      - (ex) Place diapers and beer close together



# Applications of Frequent Itemsets (2/3)

---

## ■ Related concepts

- **Items:** words
- **Baskets:** documents (e.g., Web pages, blogs, tweets)
- **Application**
  - Words appearing together in many documents may be a joint concept
    - (ex) {Brad, Angelina}

## ■ Plagiarism

- **items:** documents
- **Baskets:** sentences
  - Note that an item is “in” a basket if the sentence is in the document
- **Application**
  - Documents appearing together in many baskets share many sentences in common → plagiarism!

# Applications of Frequent Itemsets (3/3)

---

- Drug side effects
  - **Items:** drugs and side-effects
  - **Baskets:** patients
  - **Application**
    - Can detect combinations of drugs that result in particular side effects
    - But requires extension!
      - ***Absence*** of an item needs to be observed as well as presence

# Association Rules

---

- Represented as if-then rules about the contents of baskets
- The form of an association rule

$$I \rightarrow j$$

- $I = \{i_1, i_2, \dots, i_k\}$ : a set of items (i.e., itemset)
- $j$ : an item

- Implication

- If all of the items in  $I$  appear in some basket, then  $j$  is “likely” to appear in that basket as well

# The Confidence of an Association Rule

---

- Formalizes the notion of “likely”
- Definition
  - Let  $\text{conf}(I \rightarrow j)$  be the confidence of the association rule  $I \rightarrow j$ 
    - $I = \{i_1, i_2, \dots, i_k\}$ : an itemset
  - Let  $\text{support}(I)$  be the support for an itemset  $I$
  - Then,  $\text{conf}(I \rightarrow j)$  is defined as follows:

$$\text{conf}(I \rightarrow j) = \frac{\text{support}(I \cup \{j\})}{\text{support}(I)}$$

- That is, the confidence of  $I \rightarrow j$  is the probability of  $j$  given  $I$

# The Interest of an Association Rule

---

- Not all high-confidence rules are interesting

- If everyone purchases *milk*,  $X \rightarrow \text{milk}$  will have 100% confidence for any  $X$
- A rule  $I \rightarrow j$  is more valuable if it reflects a true relationship
  - $I$  somehow affects  $j$

- Definition

- The interest of a rule  $I \rightarrow j$ ,  $\text{interest}(I \rightarrow j)$ , is defined as follows:

$$\text{interest}(I \rightarrow j) = \text{conf}(I \rightarrow j) - \frac{\text{support}(\{j\})}{\# \text{ of baskets}}$$

- $\text{interest}(I \rightarrow j) = 0 \rightarrow I$  has no influence on  $j$
- $\text{interest}(I \rightarrow j) > 0 \rightarrow$  the presence of  $I$  causes the presence of  $j$
- $\text{interest}(I \rightarrow j) < 0 \rightarrow$  the presence of  $I$  discourages the presence of  $j$

# (Ex) Confidence and Interest

---

$$B_1 = \{m, c, b\}$$

$$B_2 = \{m, p, j\}$$

$$B_3 = \{m, b\}$$

$$B_4 = \{c, j\}$$

$$B_5 = \{m, p, b\}$$

$$B_6 = \{m, c, b, j\}$$

$$B_7 = \{c, b, j\}$$

$$B_8 = \{b, c\}$$

- Consider the association rule  $\{m, b\} \rightarrow c$ 
  - $\text{conf}(\{m, b\} \rightarrow c) = 2/4 = 0.5$
  - $\text{interest}(\{m, b\} \rightarrow c) = 0.5 - 5/8 = 0.5 - 0.625 = -0.125$ 
    - Item  $c$  appears in 5/8 of the baskets
    - Thus, the rule is not very interesting!

# Finding Association Rules (1/2)

---

- Problem definition

- Find all association rules  $I \rightarrow j$  with  $\text{support}(I \cup \{j\}) \geq s$  and  $\text{conf}(I \rightarrow j) \geq c$ 
  - $s$  must be reasonably high (in practice, often around 1% of the baskets)
  - $c$  must also be reasonably high (perhaps 50%)

- Suppose we have found all itemsets with  $\text{support} \geq s$

- We also have the exact support for each of these itemsets

- We then can *easily* find within them all the association rules that have both high support and high confidence

# Finding Association Rules (2/2)

---

- Let  $J$  be a set of  $n$  items that is found to be frequent
  - Then there are only  $n$  possible rules, namely  $J - \{j\} \rightarrow j$  for each  $j$  in  $J$
- For every element  $j$  of  $J$ , generate a rule  $J - \{j\} \rightarrow j$ 
  - Since  $\text{support}(J) \geq s$ , it is obvious that  $\text{support}(J - \{j\} \cup \{j\}) \geq s$
  - Thus, we only need to check if  $\text{conf}(J - \{j\} \rightarrow j) \geq c$
  - We can easily compute  $\text{conf}(J - \{j\} \rightarrow j) = \text{support}(J) / \text{support}(J - \{j\})$ 
    - Because  $J$  is frequent,  $J - \{j\}$  must be at least as frequent
    - Thus we already have  $\text{support}(J)$  and  $\text{support}(J - \{j\})$
  - Output the rule  $J - \{j\} \rightarrow j$  if  $\text{conf}(J - \{j\} \rightarrow j) \geq c$
- The **hard** part is finding all frequent itemsets!  $\rightarrow$  explained next



# (Ex) Finding Association Rules

---

$$B_1 = \{m, c, b\}$$

$$B_2 = \{m, p, j\}$$

$$B_3 = \{m, b\}$$

$$B_4 = \{c, j\}$$

$$B_5 = \{m, p, b\}$$

$$B_6 = \{m, c, b, j\}$$

$$B_7 = \{c, b, j\}$$

$$B_8 = \{b, c\}$$

- Let support threshold  $s = 3$  and confidence threshold  $c = 0.75$

## 1. Find frequent itemsets:

–  $\{b, m\}, \{b, c\}, \{c, m\}, \{c, j\}, \{m, c, b\}$

## 2. Find association rules:

~~$$\{b\} \rightarrow m: c = 4/6$$~~

$$\{m\} \rightarrow b: c = 4/5$$

$$\{b\} \rightarrow c: c = 5/6$$

~~$$\{c\} \rightarrow b: c = 3/5$$~~

~~$$\{b, c\} \rightarrow m: c = 3/5$$~~

$$\{b, m\} \rightarrow c: c = 3/4$$

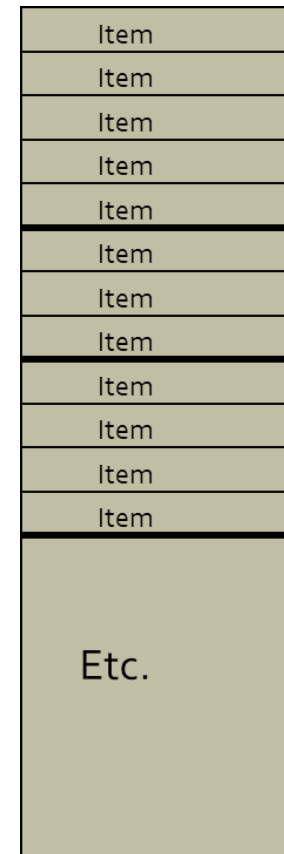
...

---

# A-Priori Algorithm

# A-Priori Algorithm

- The original algorithm to find association rules with high support and confidence
  - From which many variants have been developed
- Assumption: representation of market-basket data
  - Market-basket data is stored in a file basket-by-basket
  - Each item is represented by an *integer* code
    - (ex) {23, 456, 1001} {3, 18, 92, 145} { ... }
  - The average size of a basket is small
  - However, the size of the file of baskets is very large
    - Thus, it does not fit in main memory



# The Cost of An Association Rule Algorithm

---

- A major cost of an association rule algorithm is the time it takes to *read the baskets from disk*
  - Because the file of baskets is too large to fit in main memory
- Once a basket is in main memory, generating all the subsets of size  $k$  in main memory should take time much less than the time it took to read the basket from disk
  - (ex) if there are 20 items in a basket, there are  ${}_{20}C_2 = 190$  subsets of size 2
  - As  $k$  gets larger, the time required to generate all the subsets of size  $k$  for a basket with  $n$  items grows larger
  - However, it is usually possible to eliminate many of the items in each basket as not able to participate in a frequent itemset, so  $n$  drops as  $k$  increases

# The Cost of An Association Rule Algorithm

---

- Consequently, the work of examining each of the baskets can usually be assumed proportional to the size of the file
- In practice, association rule algorithms read the data in *passes*
  - All baskets read in turn (sequentially)
- Thus, the running time of an association rule algorithm is proportional to:
  - ***The number of passes*** taken by the algorithm  $\times$  the size of the file
- Thus, what does matter is ***the number of passes***

# Main Memory Bottleneck

---

- For many association rule algorithms, *main memory* is the critical resource
  - For example, we need to **count** the occurrence of each pair of items
    - (ex) {Bread,Milk}:10, {Milk,Beer}:8, {Coke,Beer}:4, ...
  - However, the number of pairs of items can be very large
    - For  $n$  items, the number of all possible pairs is  ${}_nC_2$
  - If we do not have enough main memory to store each of the counts, then the algorithm will **thrash**
    - Because adding 1 to a random count will most likely require us to load a page from disk
- Thus, each algorithm has a **limit** on how many items it can deal with

# Finding Frequent Pairs

---

- The hardest problem often turns out to be finding the frequent *pairs* of items  $\{i_1, i_2\}$ 
  - Why?
    - Frequent pairs are common, frequent triples are rare
    - Probability of being frequent drops exponentially with size
- Thus, let's first concentrate on pairs, then extend to larger sets
- The approach:
  - We always need to generate all the itemsets
  - But we would only like to count (keep track) of those itemsets that in the end turn out to be frequent

# Counting Pairs in Memory

---

- For  $n$  items, the number of pairs is  ${}_nC_2 = n(n-1)/2$ 
  - Suppose  $n = 10^5$  and counts are 4 byte integers
  - Naively,  $4 \times 10^5(10^5 - 1)/2 \approx 2 \times 10^{10}$  bytes = 20 GB of memory needed
- How can we reduce the amount of memory required to count?
  - Triangular matrix method
    - Store counts in a one-dimensional triangular array
  - Triple method
    - Store counts as triples  $[i, j, c]$ , meaning the count of pair  $\{i, j\}$  ( $i < j$ ) is  $c$



# Triangular Matrix Method

---

- We could use a two-dimensional array  $a$ 
  - We store the count of a pair  $\{i, j\}$  in  $a[i, j]$
  - However, this strategy makes half the array useless because  $\{i, j\} = \{j, i\}$
- A more space-efficient way
  - Use a ***one-dimensional triangular array***
  - We store in  $a[k]$  the count of a pair  $\{i, j\}$  ( $1 \leq i < j \leq n$ ), where

$$k = (i - 1) \left( n - \frac{i}{2} \right) + j - i$$

- In this layout, the pairs are stored in lexicographic order
  - $\{1, 2\}, \{1, 3\}, \dots, \{1, n\}, \{2, 3\}, \{2, 4\}, \dots, \{2, n\}, \{3, 4\}, \{3, 5\}, \dots, \{3, n\},$   
 $\dots, \{n-2, n-1\}, \{n-2, n\}, \dots, \{n-1, n\}$


# Triple Method

---


- The triangular matrix method requires us to store the count of a pair even when the count is 0
  - Inefficient when the counts of most pairs are 0
- Thus, we store triples  $[i, j, c]$  only for pairs  $\{i, j\}$  ( $i < j$ ) with  $c > 0$ 
  - Also, we also use a hash table with  $i$  and  $j$  as the search key to find a triple  $[i, j, c]$  quickly
- Comparisons
  - Triangular matrix method: uses **4 bytes** per pair (but for all pairs)
  - Triple method: uses **4+4+4=12 bytes** per pair (but only for pairs with  $c > 0$ )
  - Therefore, the triangular matrix will be better if at least  $1/3$  of  ${}_nC_2$  possible pairs actually appear in the basket data

# Comparing the Two Approaches

---



**Triangular Matrix**



**Triples**

# Monotonicity of Itemsets

---

- If we have too many items so the pairs do not fit into memory, can we do better?
- The key observation driving much of the effectiveness of the algorithms we shall discuss

**If a set  $I$  of Items is frequent, then so is every subset of  $I$**

- **Proof**
  - For  $J \subseteq I$ , every basket that contains all the items in  $I$  surely contains all the items in  $J$
  - Thus, the count for  $J$  must be at least as great as the count for  $I$
  - Consequently, if the count for  $I$  is at least  $s$ , then count for  $J$  is at least  $s$

# Maximal Frequent Itemsets

---

- Monotonicity offers us a way to compact the information about frequent itemsets
- We say an itemset is *maximal* if no superset is frequent
  - (ex)  $\text{support}(\{a\}) > s$  and  $\text{support}(\{a, b\}) > s \rightarrow \{a\}$  is *not* maximal
- Why useful?
  - If we list only the maximal itemsets, then we know that *all* of their subsets are frequent
  - No set that is not a subset of some maximal itemset can be frequent
  - Thus, we can reduce the number of output itemsets
    - (ex)  $\{a\}, \{b\}, \{a, b\} \rightarrow \{a, b\}$

# (Ex) Maximal Frequent Itemsets

---

	Support	Maximal (s=3)	
{a}	4	No	
{b}	5	No	Frequent, but superset {b, c} also frequent
{c}	3	No	
{a, b}	4	Yes	
{a, c}	2	No	Frequent, and its only superset, {a, b, c} is not frequent
{b, c}	3	Yes	
{a, b, c}	2	No	

# Tyranny of Counting Pairs

---

- In practice, the most main memory is required for determining the frequent *pairs*
  - The number of items  $n$  is rarely so large so we can count them in memory
  - The number of pairs can be very large
  - The number of *frequent* triples, quadruples, and larger sets are *rare*
    - # of frequent pairs > # of frequent triples > # of frequent quadruples > ...
- Thus, we first concentrate on algorithms for computing *frequent pairs*, then extend to larger sets
  - There are many more triples than pairs
  - It is the job of the A-Priori algorithm and related algorithms to avoid counting many triples or larger sets, and they are effective in doing so

# The A-Priori Algorithm

---

- Naïve approach (if we have enough memory)
  - Read the file of baskets in a single pass
  - For each basket of  $n$  items, generate all the  ${}_nC_2$  pairs by a double loop
  - Each time we generate a pair, add 1 to its count
  - At the end, examine all pairs and output those pairs whose count  $> s$
  - However, this approach fails if there are too many pairs of items to count them all in main memory
- A-Priori algorithm
  - Designed to reduce the number of pairs that must be counted
  - To do so, it performs **two** passes over data, rather than one pass



# The First Pass of A-Priori

---

- We create two tables

- The first table: translates item names into integers from 1 to  $n$
- The other table: an array of counts
  - The  $i$ th array element counts the occurrences of the item numbered  $i$
  - Initially, the counts for all the items are 0

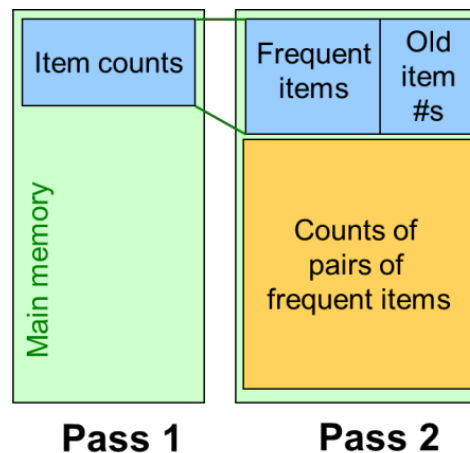
- As we read baskets

- We look at each item in the basket and translate its name into an integer
- Next, we use that integer to index into the array of counts
- We add 1 to the integer found there

- ✓ Requires only memory proportional to the number of items

# Between the Passes of A-Priori

- We examine the counts of items to determine which of them are *frequent*
- For the second pass of A-Priori, we create a new numbering from 1 to  $m$  for *just* the frequent items
  - This table is an array indexed 1 to  $n$
  - The entry for  $i$  is either 0, if item  $i$  is not frequent, or a unique integer in the range 1 to  $m$ , if item  $i$  is frequent
  - We shall refer to this table as the *frequent-item table*




# The Second Pass of A-Priori

---

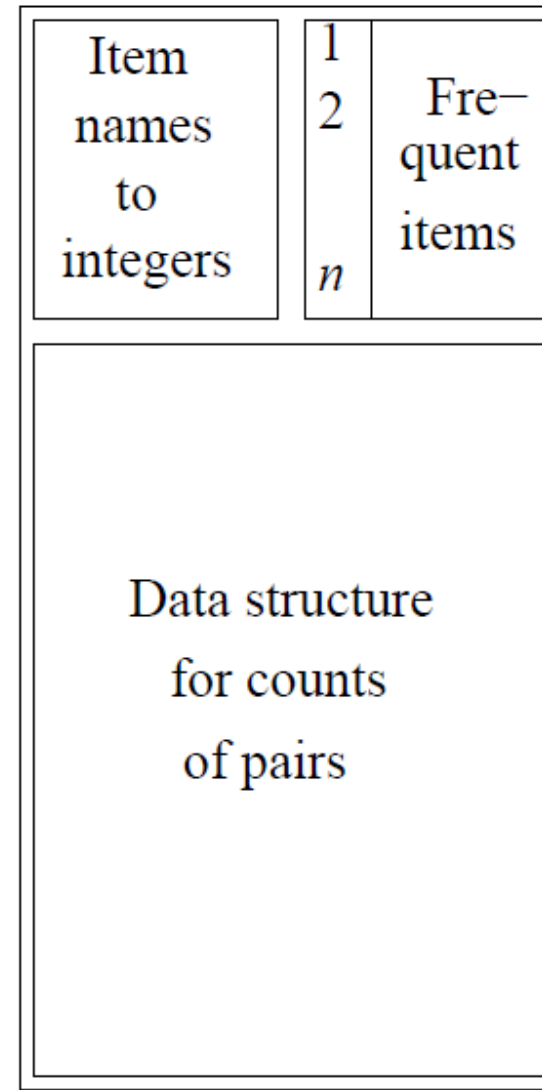
- We count all the pairs that consists of two *frequent* items
  - Recall that a pair cannot be frequent unless both its items are frequent
  - If we use the triangular-matrix method
    - The space required is  $2m^2$  bytes, rather than  $2n^2$  bytes
  - If we use the triples method
    - The renumbering of just the frequent items is necessary
- The mechanics
  - For each basket, look in the frequent-items table to find frequent items
  - In a double loop, generate all pairs of frequent items in that basket
  - For each such pair, add 1 to its count
  - Finally, examine the counts to determine which pairs are frequent
- ✓ Requires memory proportional to square of *frequent* items only

# Main Memory Use of A-Priori Algorithm

---



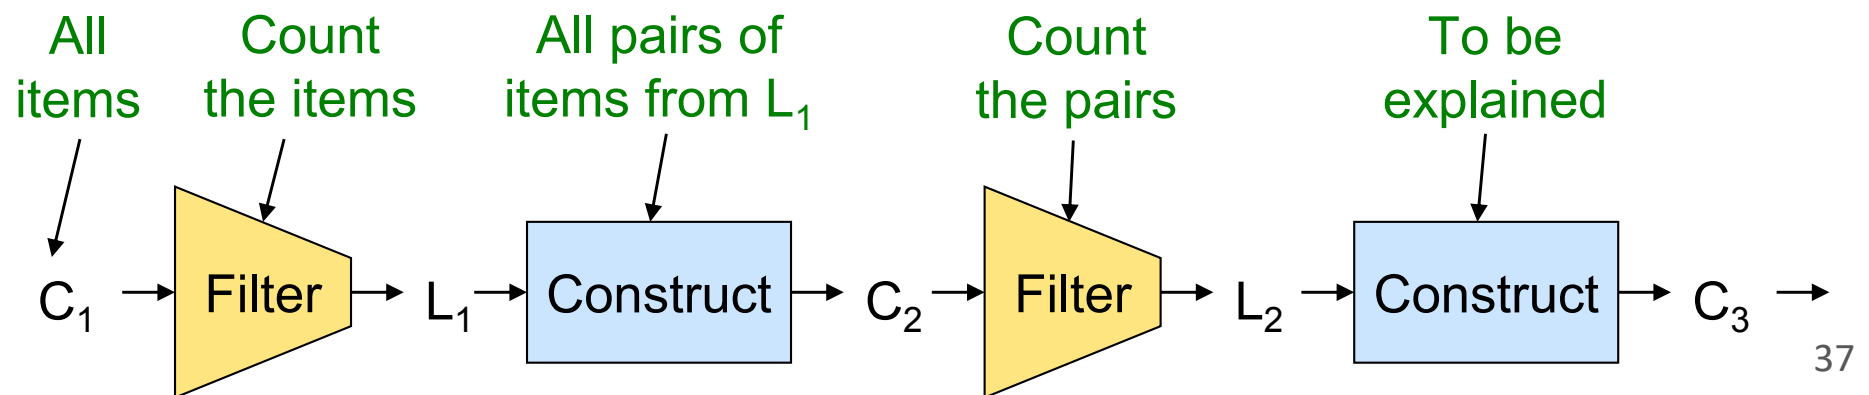
Pass 1



Pass 2

# A-Priori for All Frequent Itemsets

- We can use the *same* technique to find larger frequent itemsets without an exhaustive count of all sets
- For each size  $k = 1, 2, \dots$ , we construct two sets of itemsets
  - $C_k$  = the set of candidate itemsets of size  $k$ 
    - The itemsets we must count in order to determine whether they are frequent
  - $L_k$  = the set of truly frequent itemsets of size  $k$
  - If  $L_k = \emptyset$ , then we stop
    - By monotonicity, there can be no larger frequent itemsets



# Constructing $C_k$ from $L_{k-1}$

---

## ■ Step 1: Join

- Generate  $C_k$  by merging a pair of frequent itemsets in  $L_{k-1}$  if their  $(k-2)$  items are identical

## ■ Step 2: Prune

- Delete all itemsets  $I \in C_k$  if some  $(k-1)$ -subset of  $I$  is not in  $L_{k-1}$ 
  - Because all  $(k-1)$ -subsets of  $I$  must be frequent if  $I$  is frequent
  - In other words,  $C_k$  is all those itemsets of size  $k$ , every  $k-1$  of which is an itemset in  $L_{k-1}$

## ■ Example

- Let  $L_3 = \{\{1, 2, 3\}, \{1, 2, 4\}, \{1, 3, 4\}, \{1, 3, 5\}, \{2, 3, 4\}\}$
- After the join step,  $C_4 = \{\{1, 2, 3, 4\}, \{1, 3, 4, 5\}\}$
- After the prune step,  $C_4 = \{\{1, 2, 3, 4\}\}$  ( $\because \{1, 4, 5\}$  is not in  $L_3$ )

# Constructing $L_k$ from $C_k$

---

- Make a pass through the baskets and counting all and only the itemsets of size  $k$  that are in  $C_k$
- Those itemsets that have count at least  $s$  are in  $L_k$
- If we find  $L_k = \emptyset$ , then we stop

# (Ex) A-Priori for All Frequent Itemsets

---

- ① Let  $C_1 = \{\{b\}, \{c\}, \{j\}, \{m\}, \{n\}, \{p\}\}$
- ② Count the support of itemsets in  $C_1$
- ③ Prune non-frequent:  $L_1 = \{\{b\}, \{c\}, \{j\}, \{m\}\}$
- ④ Construct  $C_2 = \{\{b, c\}, \{b, j\}, \{b, m\}, \{c, j\}, \{c, m\}, \{j, m\}\}$
- ⑤ Count the support of itemsets in  $C_2$
- ⑥ Prune non-frequent:  $L_2 = \{\{b, m\}, \{b, c\}, \{c, m\}, \{c, j\}\}$
- ⑦ Construct  $C_3 = \{\{b, m, c\}, \{\cancel{b, c, j}\}, \{\cancel{b, m, j}\}, \{\cancel{c, m, j}\}\}$
- ⑧ Count the support of itemsets in  $C_3$
- ⑨ Prune non-frequent:  $L_3 = \{\{b, m, c\}\}$



---

# Handling Larger Datasets

# Handling Larger Datasets in Main Memory

---

- The A-Priori algorithm is fine *as long as* the counting of the candidate pairs  $C_2$  can be accomplished in memory
  - Otherwise, **thrashing** occurs
    - i.e., repeated moving of data between disk and main memory
- Several algorithms have been proposed to **cut down** on the size of candidate set  $C_2$ 
  - PCY algorithm
  - Multistage algorithm
  - Multihash algorithm

# PCY (Park-Chen-Yu) Algorithm

---

## ■ Observation

- There may be much **unused** space in main memory on the first pass
  - We do not need more than the main memory for the two tables
  - A translation table from items to integers and an array to count those integers
  - (ex) If there are  $10^6$  items, only  $12 \times 10^6$  bytes = 12 MB is used for the counts


## ■ Can we use the unused memory to **reduce** memory required in the second pass?

## ■ The idea of PCY

- Use the unused space to store some information that can **filter** infrequent pairs in the second pass
- Similar to the idea of a Bloom filter

# Main Memory Use of PCY

---



# The First Pass of PCY

---

- In addition to item counts, maintain a *hash table* with as many buckets as fit in memory
  - As we examine a basket during the first pass, we not only count each item in the basket, but generate all pairs, using a double loop
  - We hash each pair and add 1 to the bucket into which that pair hashes
  - Note that we just keep the count for each bucket, **not** the actual pairs

```
FOR (each basket) :  
    FOR (each item in the basket) :  
        add 1 to item's count;  
New in PCY { FOR (each pair of items) :  
                hash the pair to a bucket;  
                add 1 to the count for that bucket;
```

# Observation about Buckets

---

- At the end of the first pass, each bucket has a count
  - The sum of counts of all the pairs that hash to that bucket
- If the count of a bucket  $\geq s$  (i.e., the support threshold)
  - The bucket is called a **frequent** bucket
  - The pairs that hash to a frequent bucket **may or may not** be frequent
- However, if the count of a bucket  $< s$ 
  - The bucket is called an **infrequent** bucket
  - **no** pair that hashes to an infrequent bucket can be frequent
- Thus, we **only** count pairs that hash to **frequent** buckets in the second pass!

# Between The Passes of PCY

---

- The hash table is replaced by a bitmap, with 1 bit for each bucket
  - The bit is 1, if the bucket is frequent and 0, if not
- Space reduction
  - 4 byte integers for each bucket are replaced by a single bit
  - Thus, the bitmap takes only **1/32** of the space
- If most buckets are infrequent (this is what we expect)
  - The number of pairs being counted on the second pass will much ***smaller*** than the total number of pairs of frequent items
  - Thus, PCY can handle some data sets ***without*** thrashing during the second pass, while A-Priori would run out of main memory and thrash

# The Second Pass of PCY

---

- Count only those pairs  $\{i, j\}$  that meet the following conditions for being a candidate pair in  $C_2$ :
  - ① Both  $i$  and  $j$  are frequent items
  - ② The pair  $\{i, j\}$  hashes to a bucket whose bit in the bitmap is 1
    - That is,  $\{i, j\}$  hashes to a **frequent** bucket
- Both conditions are necessary for the pair to have a chance of being frequent
- It is the **second condition** that distinguishes PCY from A-Priori



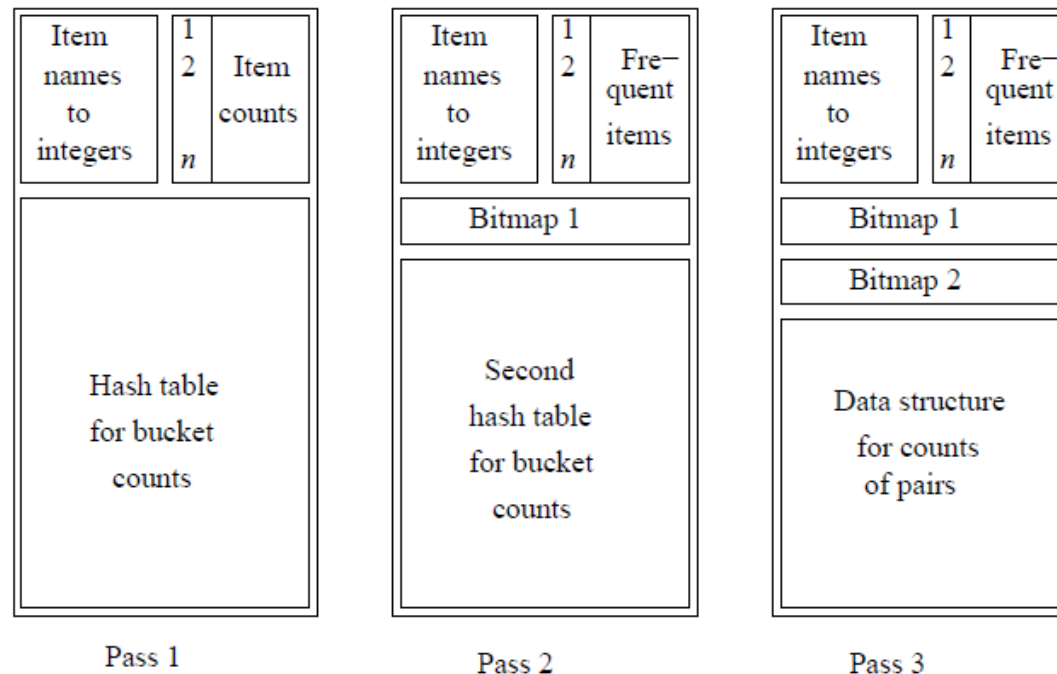
# Counting Pairs in PCY

---

- However, on the second pass of PCY, we *cannot* use the triangular matrix method to count pairs
  - Because the pairs of frequent items are placed *randomly* within the triangular matrix, we cannot compact the matrix
- Consequently, we are forced to use the *triples methods* in PCY
  - Note that A-Priori can use the both methods
  - That restrict may not matter if the fraction of pairs of frequent items that actually appear in buckets were small, but ...
- Thus, unless PCY lets us avoid counting at least **2/3** of the pairs of frequent items, we cannot gain by using PCY instead of A-Priori
  - Triangular matrix: 4 bytes per pair
  - Triple method: 12 bytes per pair

# Multistage Algorithm

- Improves PCY to reduce *further* the number of candidate pairs
  - By using *several* successive hash tables
  - However, the tradeoff is that it requires more than two passes
- Takes *more than* two passes to find the frequent pairs
  - The first pass → the second pass → the third pass



# The First and Second Passes

---

- The first pass

- The same as the first pass of PCY
- After that pass, the frequent buckets are summarized by a bitmap

- The second pass

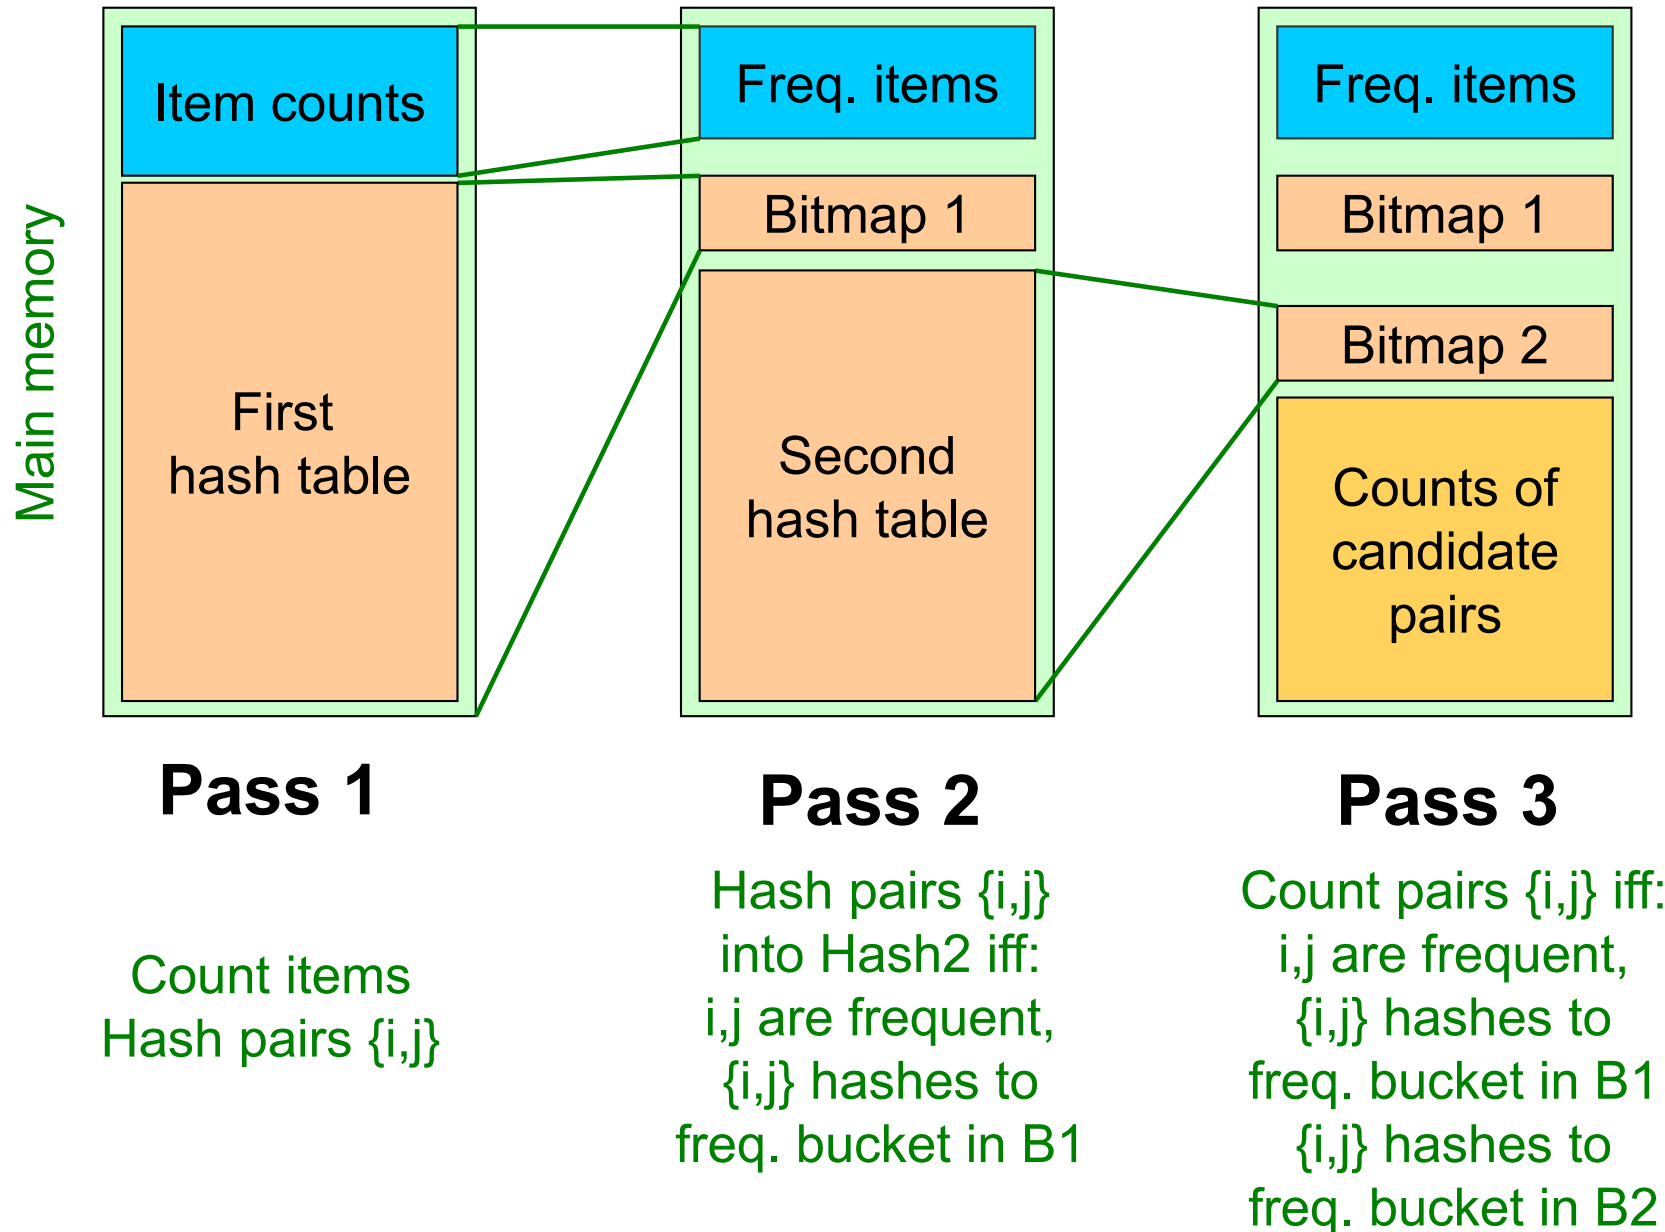
- Again go through the file of baskets
- Using **another** hash function, rehash only those pairs  $\{i, j\}$  that meet the following two criteria:
  - ① Both  $i$  and  $j$  are frequent items
  - ② The pair  $\{i, j\}$  hashed to a frequent bucket on the **first** pass
- Summarize the second hash table to a bitmap
  - As result, we expect there to be many **fewer** frequent buckets in the second hash table than in the first

# The Third Pass

---

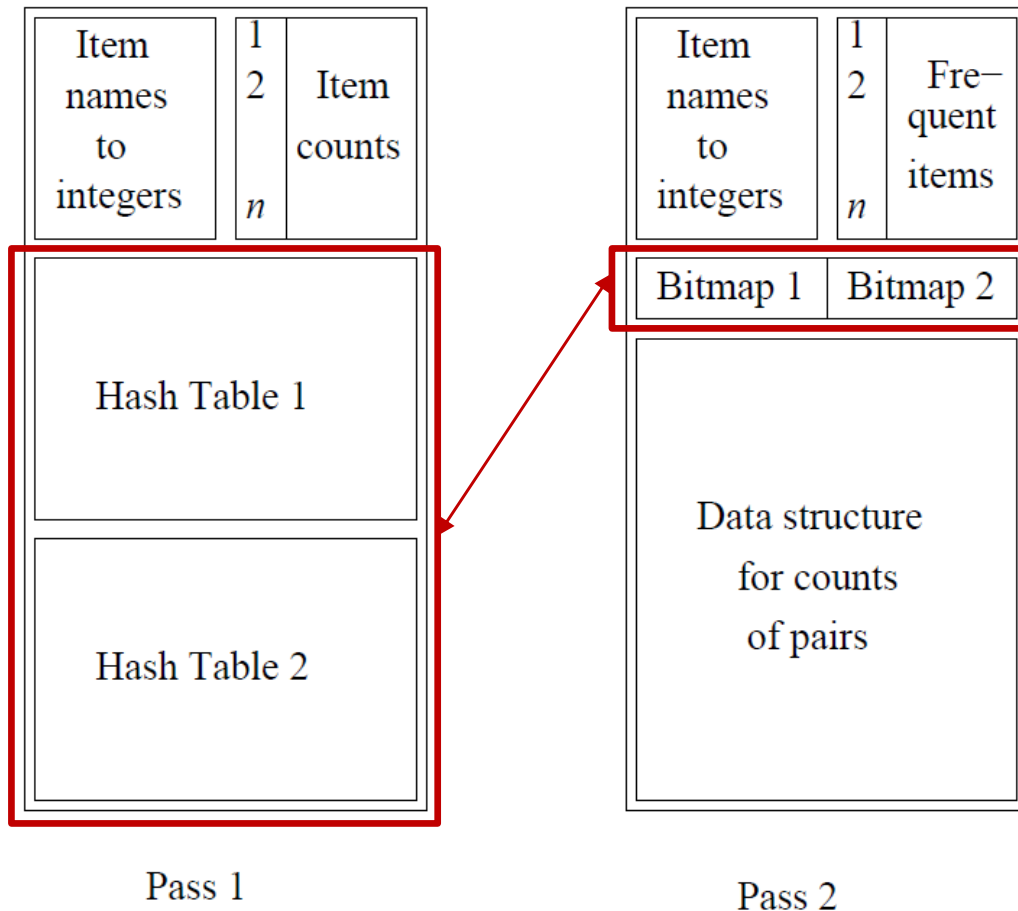
- Count only those pairs  $\{i, j\}$  that meet the following conditions for being a candidate pair in  $C_2$ :
  - ① Both  $i$  and  $j$  are frequent items
  - ②  $\{i, j\}$  hashes to a frequent bucket in the **first** hash table
  - ③  $\{i, j\}$  hashes to a frequent bucket in the **second** hash table
- The **third condition** is the distinction between Multistage and PCY
- Important points
  - The two hash functions have to be **independent**
  - We need to check **both** hashes on the third pass
    - If not, we would end up counting pairs of frequent items that hashed first to an infrequent bucket but happened to hash second to a frequent bucket

# Main Memory Use of Multistage



# Multihash Algorithm

- Key idea
  - Get the benefit of the extra passes of Multistage in a *single* pass
  - Use *several independent* hash tables on the first pass



# The Danger of Using Two Hash Tables

---

- Each hash table has *half* as many buckets as the one large hash table of PCY
  - Thus, the average count of a bucket *doubles*
  - We have to sure most buckets will still not reach count  $s$

1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---

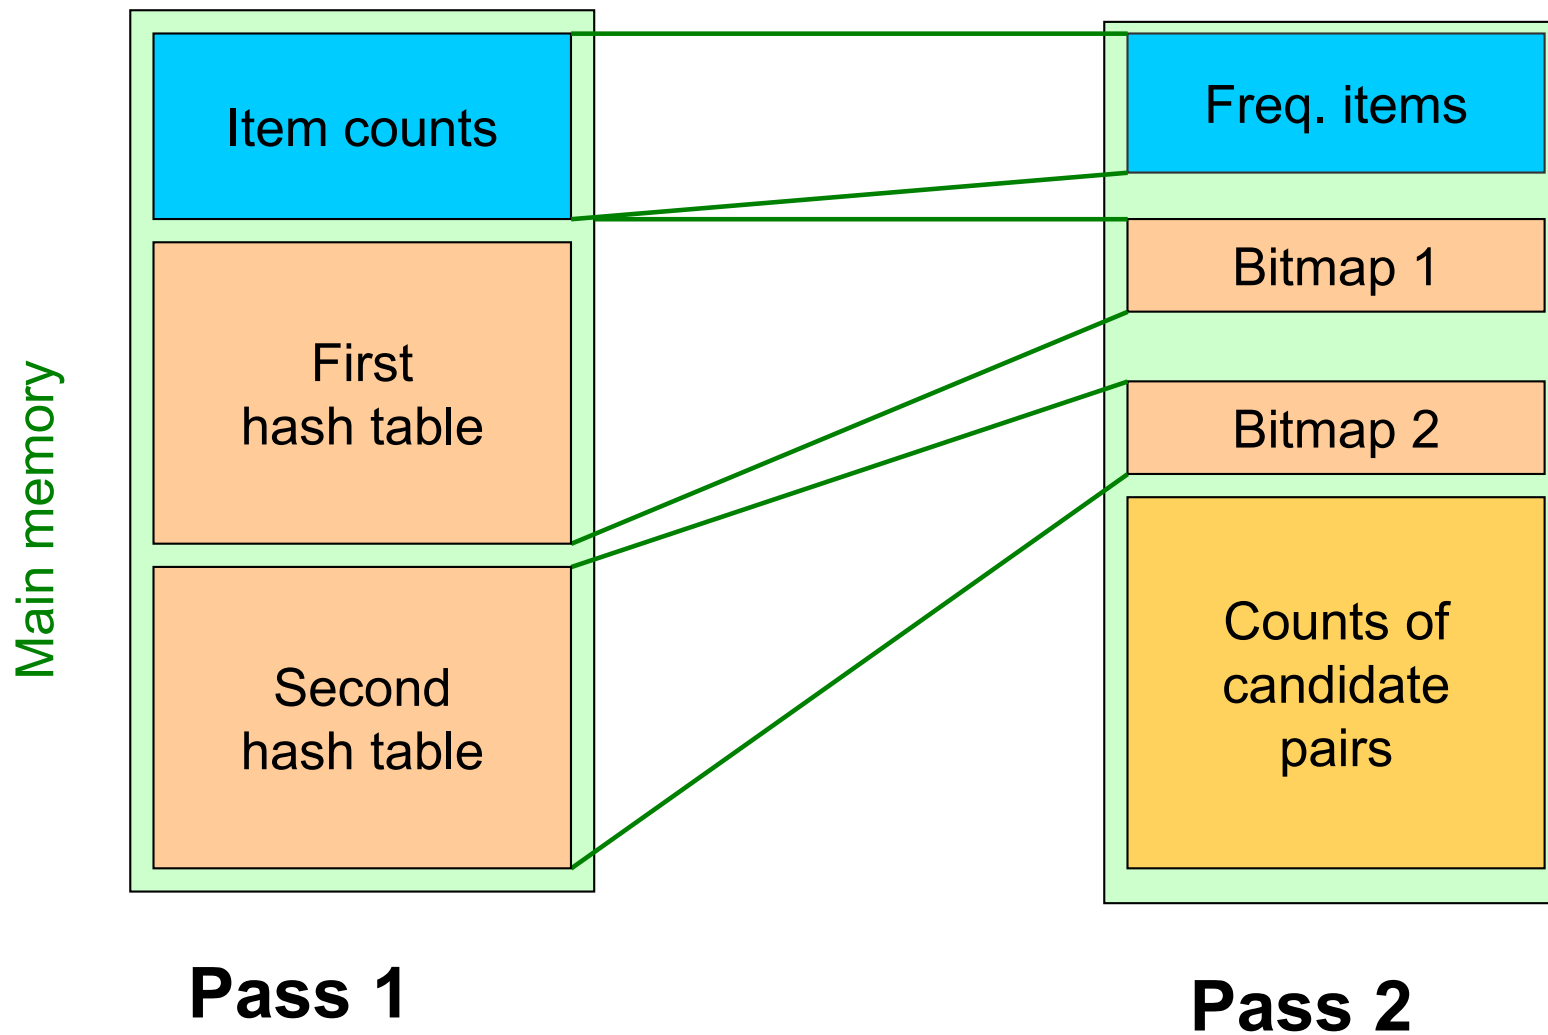
A single large hash table

2	2	2	2
2	2	2	2

Two small hash tables

- If the average count of a bucket for PCY is much lower than  $s$ 
  - We can operate two half-sized hash tables and still expect most of the buckets of both hash tables to be infrequent
  - Thus, in this situation we might well choose the multihash approach
- If so, we can get a benefit like multistage, but in only *two* passes

# Main Memory Use of Multihash





# The Second Pass of Multihash

---

- Each hash table is converted to a bitmap, as usual
- Note
  - The two bitmaps for the two hash functions occupy exactly as much space as a single bitmap would for the second pass of PCY
- The conditions for a pair  $\{i, j\}$  to be in  $C_2$ , and thus to require a count on the second pass, are the same as for the third pass of Multistage
  - ① Both  $i$  and  $j$  are frequent items
  - ②  $\{i, j\}$  hashes to a frequent bucket in the **first** hash table
  - ③  $\{i, j\}$  hashes to a frequent bucket in the **second** hash table

# Extensions of PCY

---

- Either Multistage or Multihash can use *more than* two hash functions
- Multistage (with more than two hash functions)
  - There is a point of diminishing returns, since the bit-vectors eventually consume all of main memory
- Multihash (with more than two hash functions)
  - The bit-vectors occupy exactly what one PCY bitmap does, but too many hash functions makes all counts  $\geq s$

---

# Limited-Pass Algorithm

# Limited-Pass Algorithms

---

- The algorithms for frequent itemsets discussed so far use one pass for *each* size of itemset
  - i.e., a total of  $k$  passes to find frequent itemsets of size 1, 2, ...,  $k$
  - Can we use fewer passes?
- However, there are many applications where it is *not* essential to discover every frequent itemset
  - In many cases, it is quit sufficient to find most but not all of the frequent itemsets (e.g., a supermarket application)
- We explore some algorithms that uses *at most two* passes
  - But we may miss some frequent itemsets
  - (ex) Random sampling, SON (Savasere, Omiecinski, and Navathe)

# Random Sampling (1/2)

---

- Key idea
  - Instead of using the entire file of baskets, pick a random subset of the baskets and pretend it is the entire dataset
- We ***must*** adjust the support threshold  $s$  to  $p \cdot s$  to reflect the smaller number of baskets
  - $p$ : the sampling rate
  - (ex) if we choose a sample of 1% of the baskets, we adjust  $s$  to  $s/100$
- How can we pick the sample?
  - For each basket, select that basket with some fixed probability  $p$
  - (ex) if there are  $m$  baskets in the entire file, we shall sample  $p \cdot m$  baskets

# Random Sampling (2/2)

---

- Having selected a sample, we store these baskets in main memory and execute one of the algorithms *in main memory*
  - A-Priori, PCY, Multistage, or Multihash
- Although the algorithm must run passes over the main-memory sample, there are **no** disk accesses needed to read the sample
  - Since it resides in main memory
  - Therefore, **only two** passes are required
- Of course the algorithm will fail if it can't be run in main memory
  - An option is to read the sample from disk for each pass
  - Since the sample is much smaller than the full dataset, we still avoid most of the disk I/O

# Avoiding Errors in Sampling Algorithms

---

- In random sampling, there can be *errors*
  - False negative
    - An itemset that is frequent in the whole but *not* in the sample
  - False positive
    - An itemset that is frequent in the sample but *not* in the whole
- If the sample is large enough, there are unlikely to be serious errors
  - An itemset whose support is much *larger* than  $s$  will almost surely be identified from a random sample
  - An itemset whose support is much *less* than  $s$  is unlikely to appear frequent in the sample

# How Can We Reduce Such Errors?

---

- We can eliminate *false positives*
  - Make a pass through the full dataset and count all the itemsets that were identified as frequent in the sample
  - Retain as frequent only those itemsets that were frequent in the sample and also frequent in the whole
  - But requires more passes
- We cannot eliminate *false negatives* but can reduce their number
  - Use a support threshold **smaller** than  $p \cdot s$  (e.g.,  $0.9p \cdot s$ ) to catch more truly frequent itemsets
  - We shall identify, as having support at least  $0.9p \cdot s$  in the sample, almost all those itemsets that have support at least  $s$  in the whole
  - But requires more space



# SON Algorithm

---

- Improves the random sampling to avoid **both** false negatives and false positives, at the cost of making **two** passes
- First pass
  - Divide the input file into **chunks**
  - Treat each chunk as a sample, and find all frequent items in that chunk
    - We use  $p \cdot s$  as the threshold, if each chunk is fraction  $p$  of the whole file
  - Store on disk all the frequent itemsets found for each chunk
  - These frequent itemsets become the **candidate** itemsets
- Second pass
  - Count all the candidate itemsets and select those that have support at least  $s$  as the frequent itemsets

# SON Algorithm: Key Idea

---

- In the SON algorithm, an itemset becomes a candidate if it is frequent in **any** one or more chunks
- Key “monotonicity” idea
  - An itemset **cannot** be frequent in the whole unless it is frequent in at least one chunk
  - Proof
    - If an itemset is not frequent in any chunk, then its support is less than  $p \cdot s$  in each chunk
    - Since the number of chunks is  $1/p$ , the total support of that items is less than  $(1/p)p \cdot s = s$
- Therefore, there are **no** false negatives
  - Every itemset frequent in the whole is frequent in at least one chunk

# SON Algorithm and MapReduce

---

- SON algorithm lends itself *well* to a parallel-computing environment
- Distributed version
  - ① Distribute baskets among many nodes
  - ② Compute frequent itemsets at each node
  - ③ Distribute the candidates to all nodes
  - ④ Count the support for each candidate at each node
  - ⑤ Finally sum those supports to get the support for each candidate
- SON algorithm can be easily expressed as *MapReduce* jobs
  - Of course, this process does not have to be implemented in MapReduce

# SON: MapReduce Version (1/2)

---

- **Phase 1:** Find candidates itemsets
  
- **Map task**
  - Input: a subset of the baskets
  - Output: a set of key-value pairs  $(F, 1)$ 
    - $F$ : a frequent itemset from the subset
  
- **Reduce task**
  - Input: a set of key-values pairs  $(F, [1, 1, \dots])$ 
    - $F$ : a frequent itemset identified from some subset
  - Output: a set of key-value pairs  $(F, 1)$ 
    - $F$ : a candidate itemset

# SON: MapReduce Version (2/2)

---

- **Phase 2: Find true frequent itemsets**
- **Map task**
  - Input: a subset of the candidate itemsets, a subset of the baskets
  - Output: a set of key-value pairs  $(C, v)$ 
    - $C$ : a candidate itemset
    - $v$ : the count for  $C$  among the subset of the baskets
- **Reduce task**
  - Input: a set of key-values pairs  $(C, [v_1, v_2, \dots])$ 
    - $C$ : a candidate itemset
  - Output: a set of key-value pairs  $(C, sum)$ 
    - $C$ : a frequent itemset whose  $sum$  is at least  $s$
    - $sum$ : the sum of values  $v_1, v_2, \dots$